

Reading and Writing Data to Kinesis

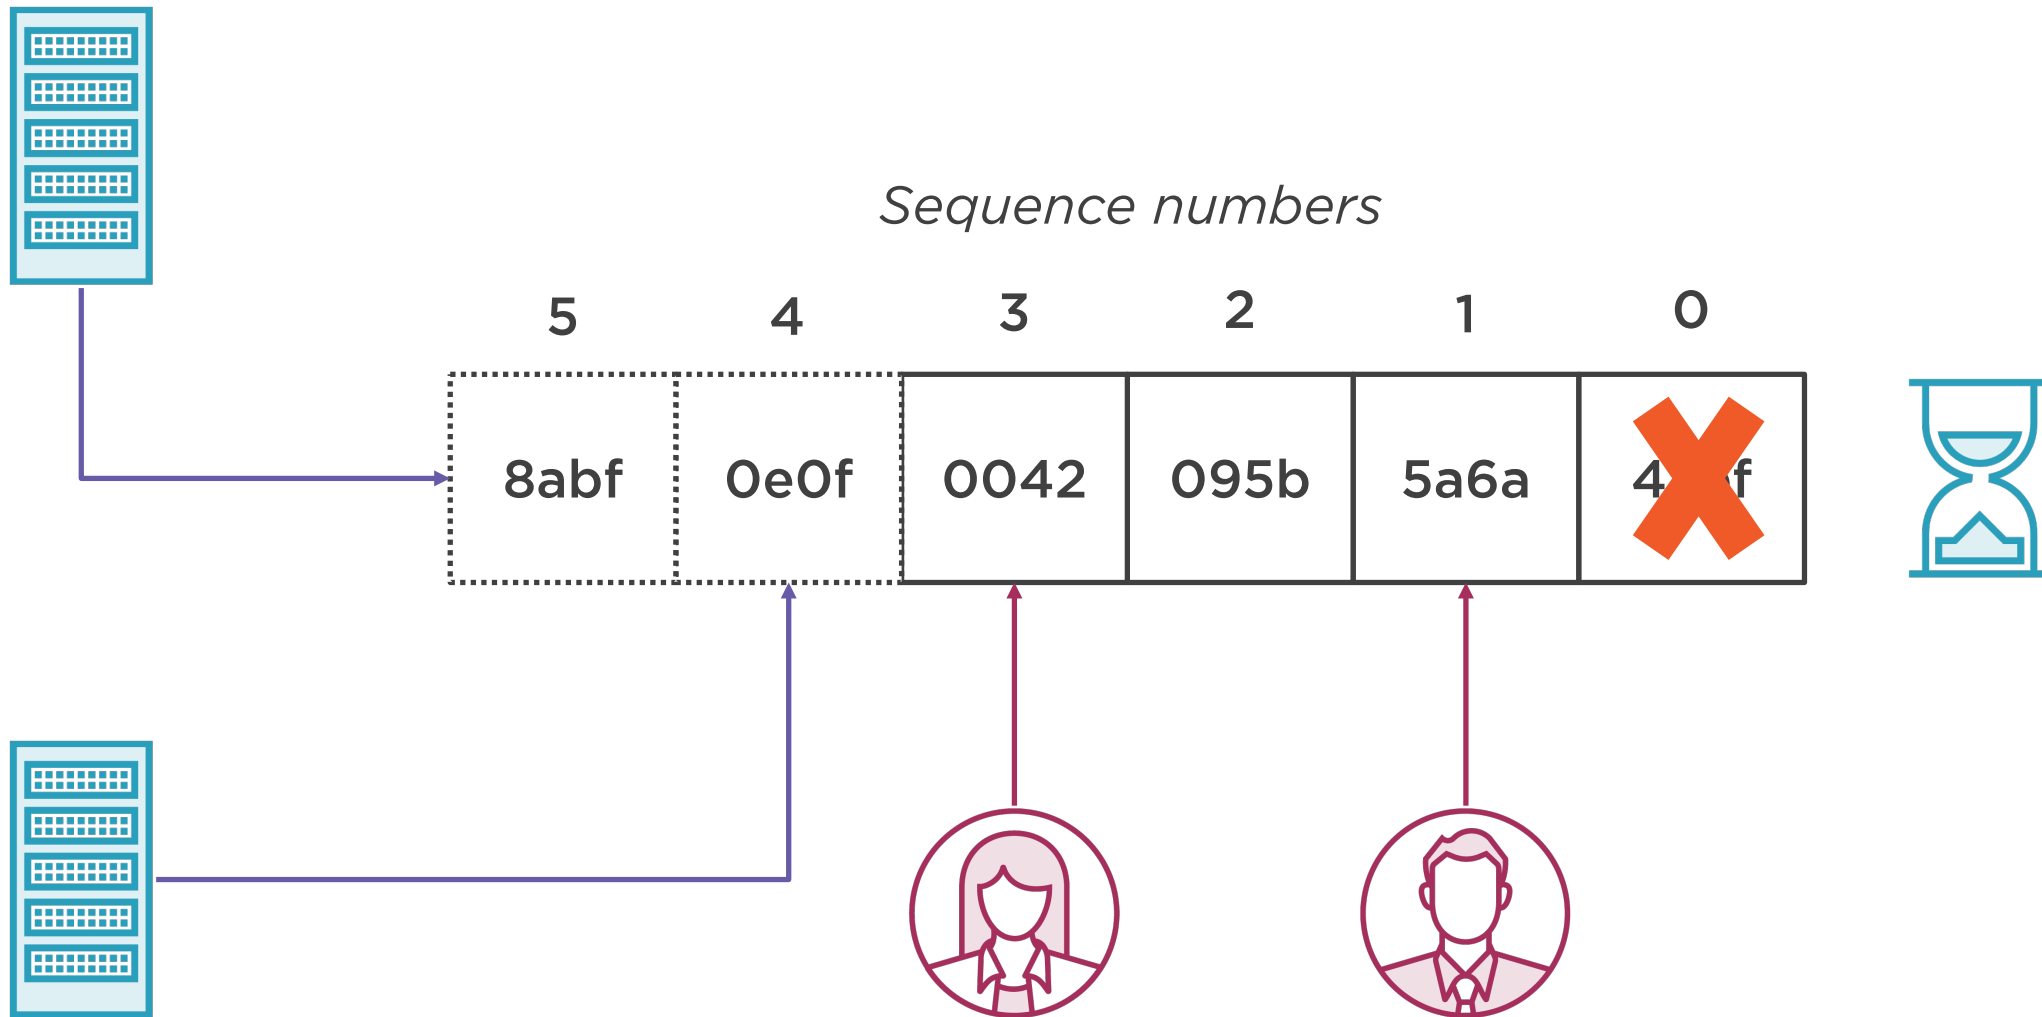


Ivan Mushketyk

@mushketyk brewing.codes



Kinesis Stream



Module Overview



Learn Kinesis API

Implement Kinesis producer


Implement Kinesis consumer

- Two different APIs to do this

Scale stream up and down



What Are We Going to Implement



Kinesis API

Stream

CreateStream
DeleteStream
ListStreams
DescribeStream

Manage Shards

MergeShards
SplitShard

Manage Stream

IncreateStreamRetentionPeriod
EnableEnhancedMonitoring
StartStreamEncryption

Write

PutRecord
PutRecords

Read

GetShardIterator
GetRecord

Fan-out Consumers

RegisterStreamConsumer
DeregisterStreamConsumer
ListStreamConsumers



```
{  
  "StreamName": "tweetsStream"  
  "PartitionKey": "NoSQL",  
  "Data": "dba7..768=",  
}
```

Writing Data to Kinesis

Writing a Tweet message to the "tweetsStream" stream with the "NoSQL" partition key



Reading Data from Twitter



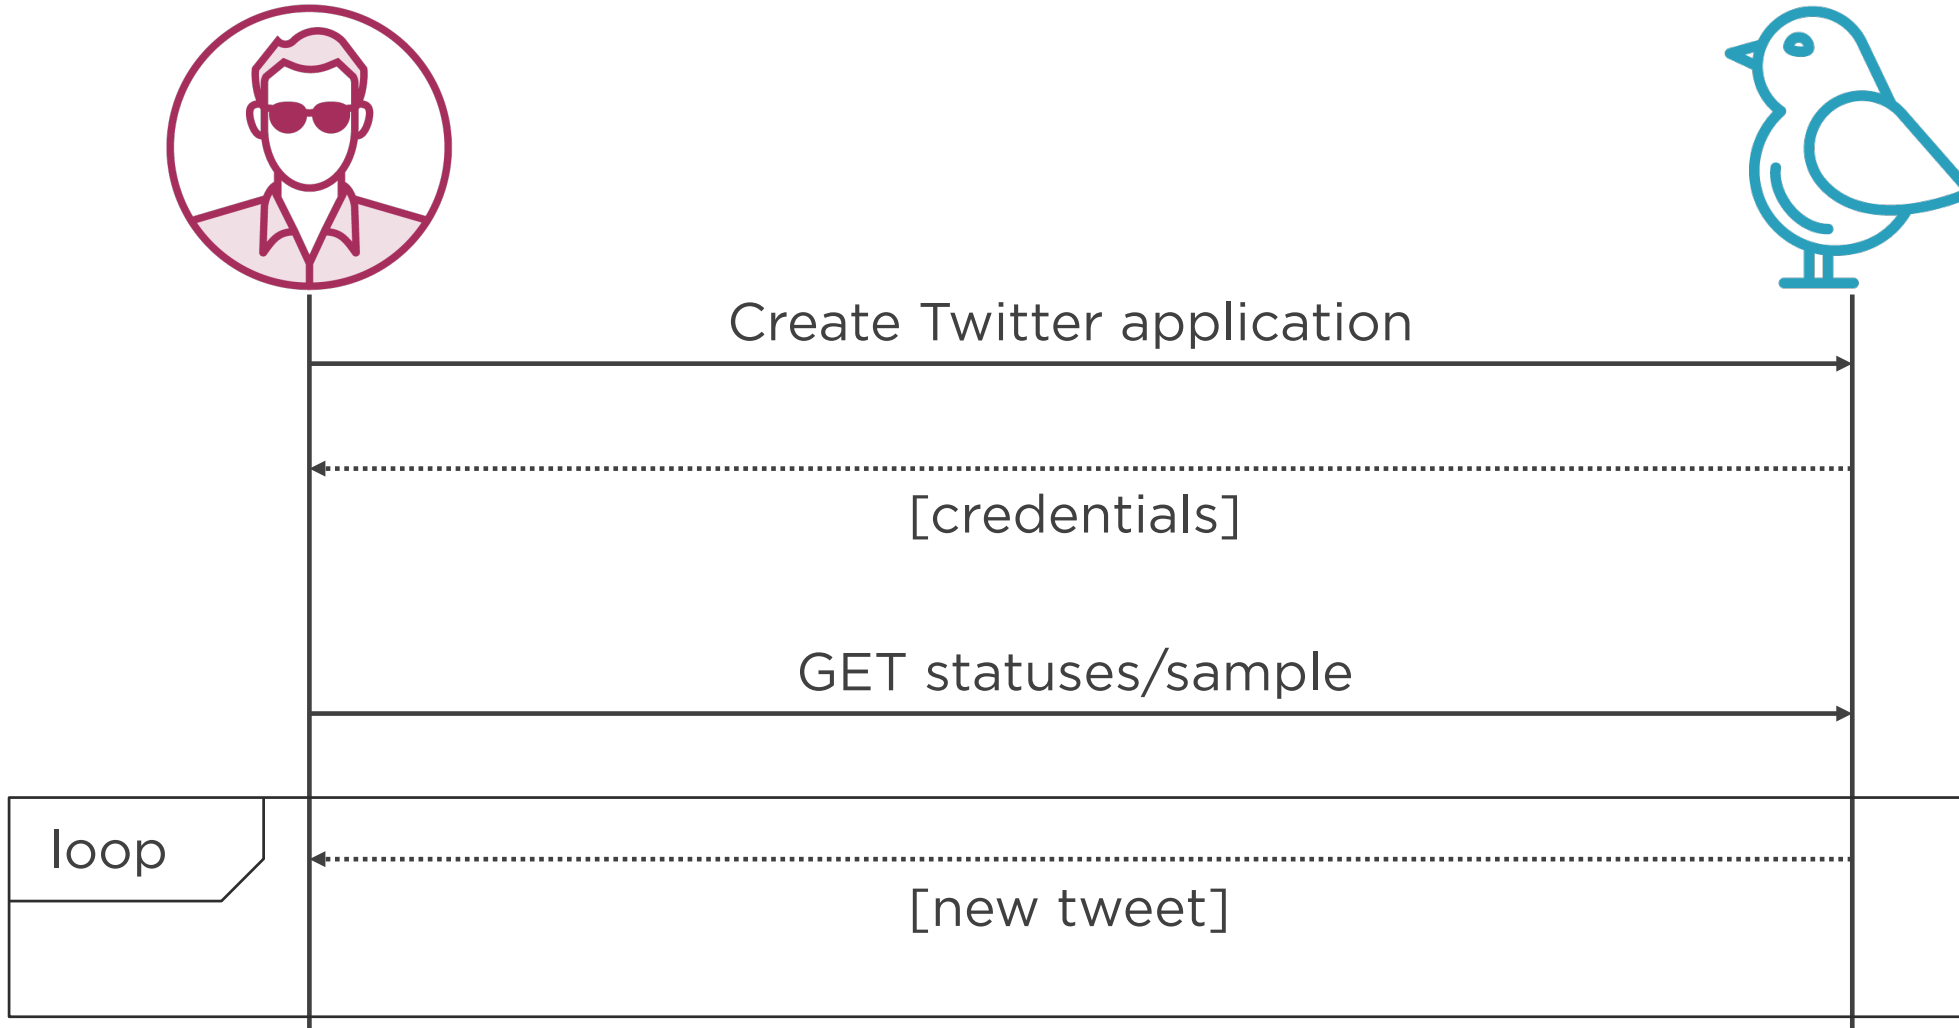
Read tweets from Twitter

- Create a Twitter application
- Use twitter4j

Main data source

Sample of tweets

Reading Data from Twitter



Demo



Create Twitter application

Read a stream of tweets



Demo



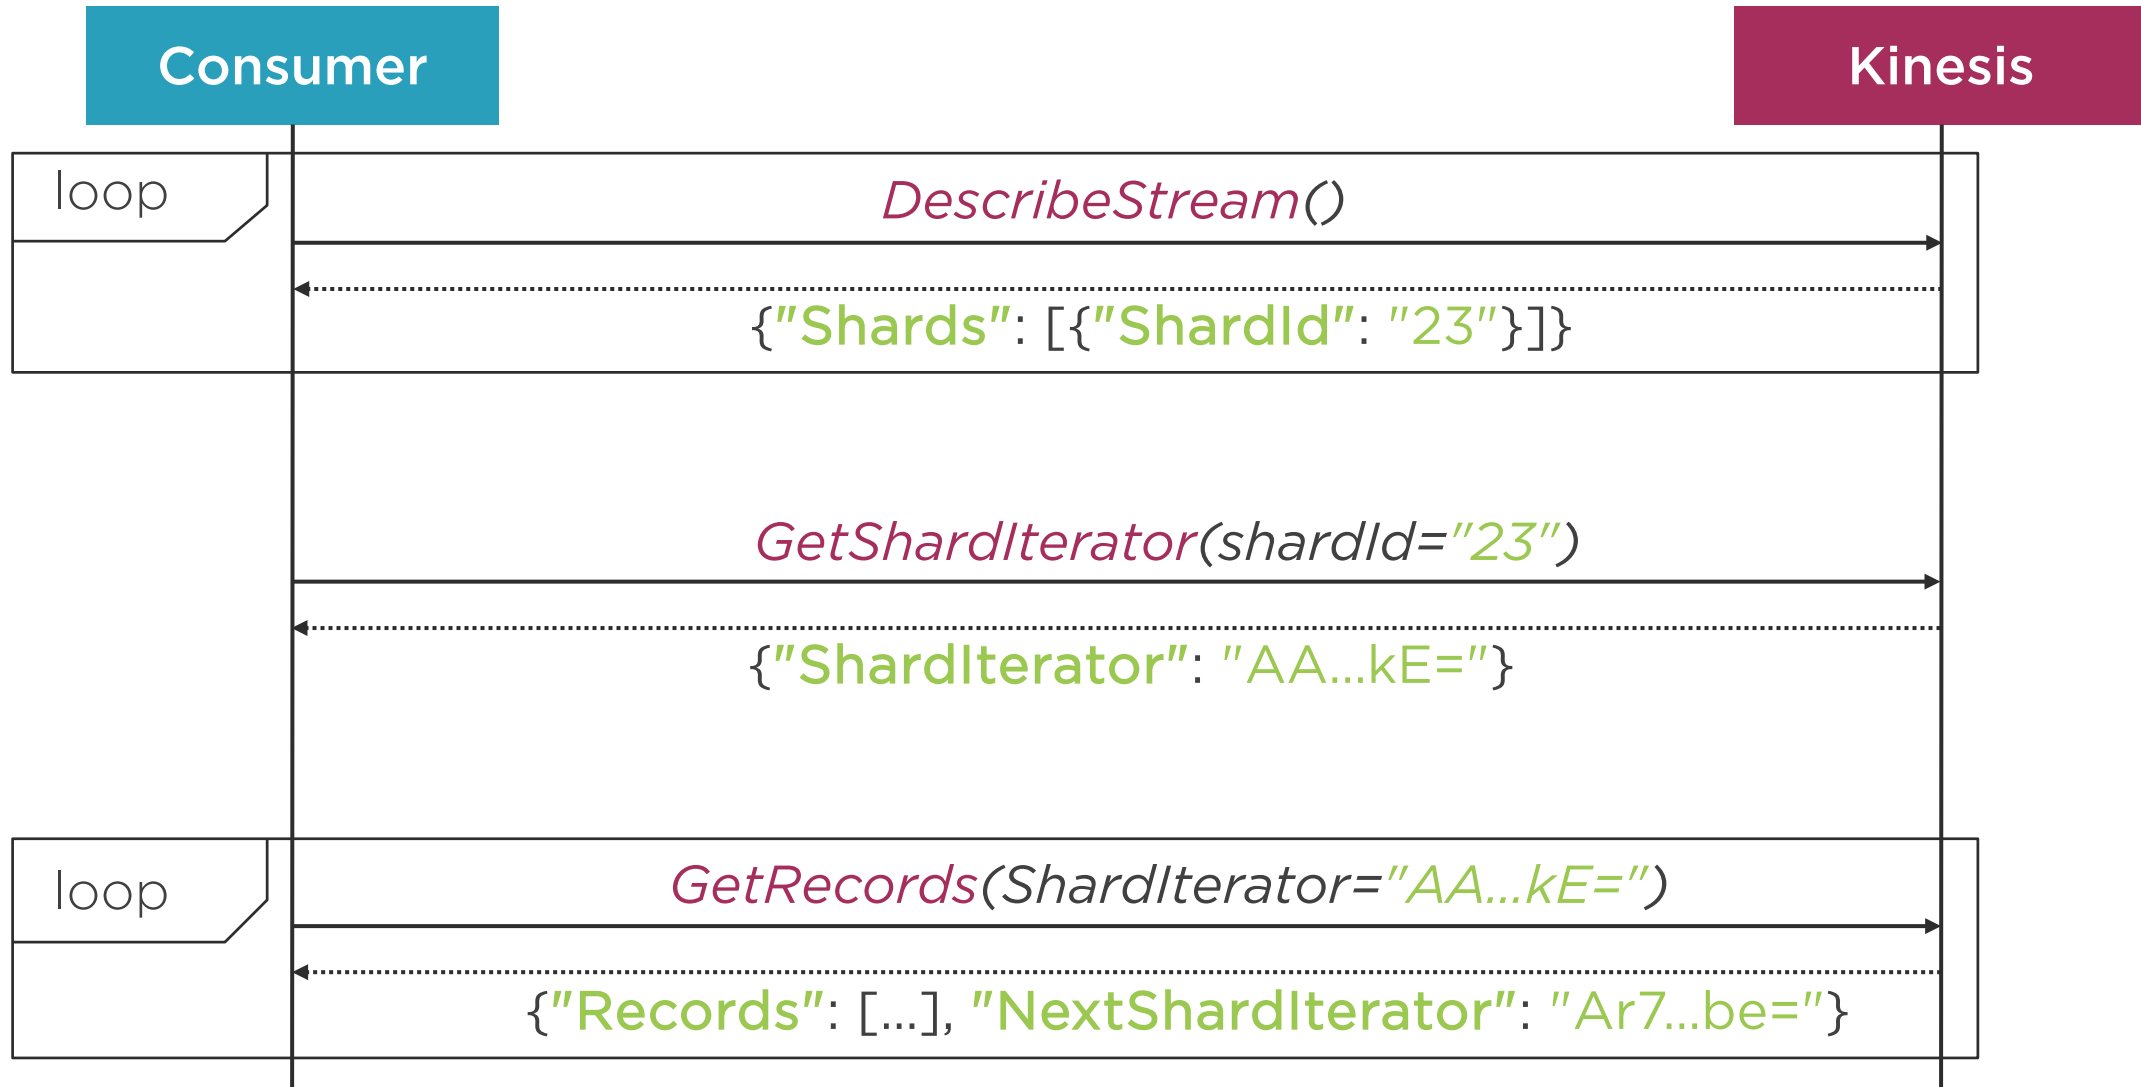
Implement Kinesis producer

Write tweets data to Kinesis

Will read it from Kinesis later



Reading Data from Kinesis



Demo



Read tweets records from Kinesis

Write tweets to console



Consumer Implementation



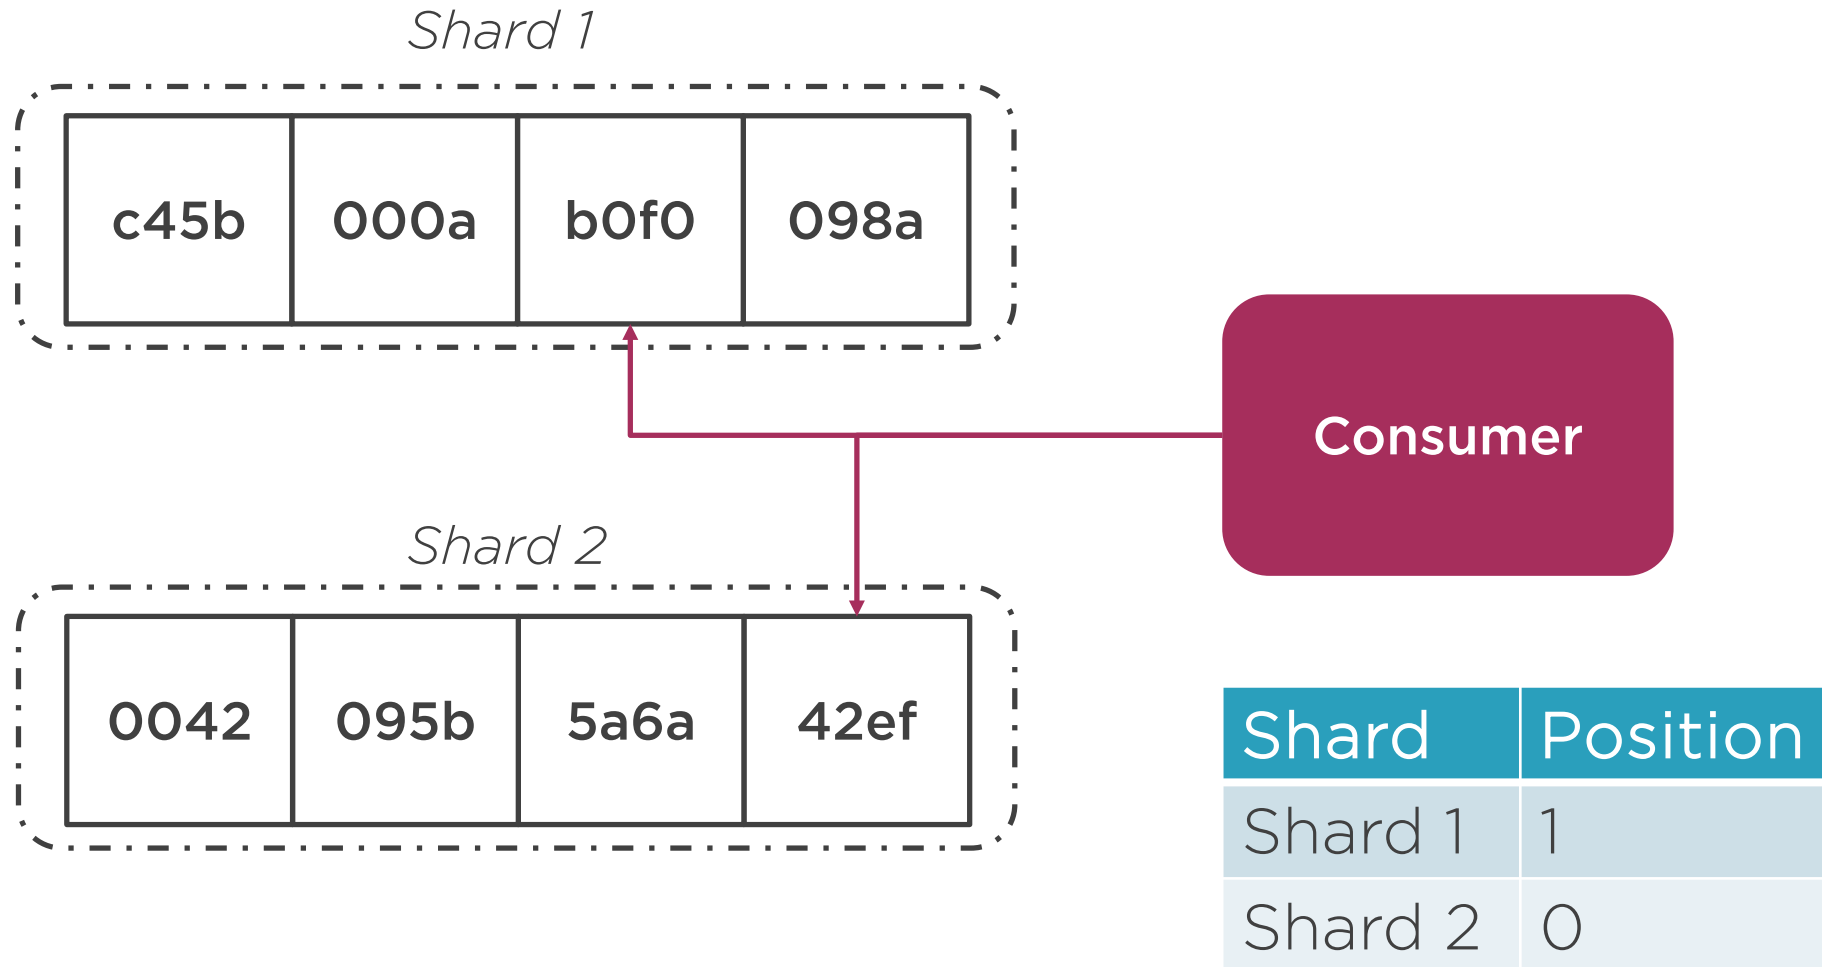
Correct consumer implementation

- Process all shards
- Assign shards to different workers
- Track progress in each shard
- Re-balance work if number of shard changes

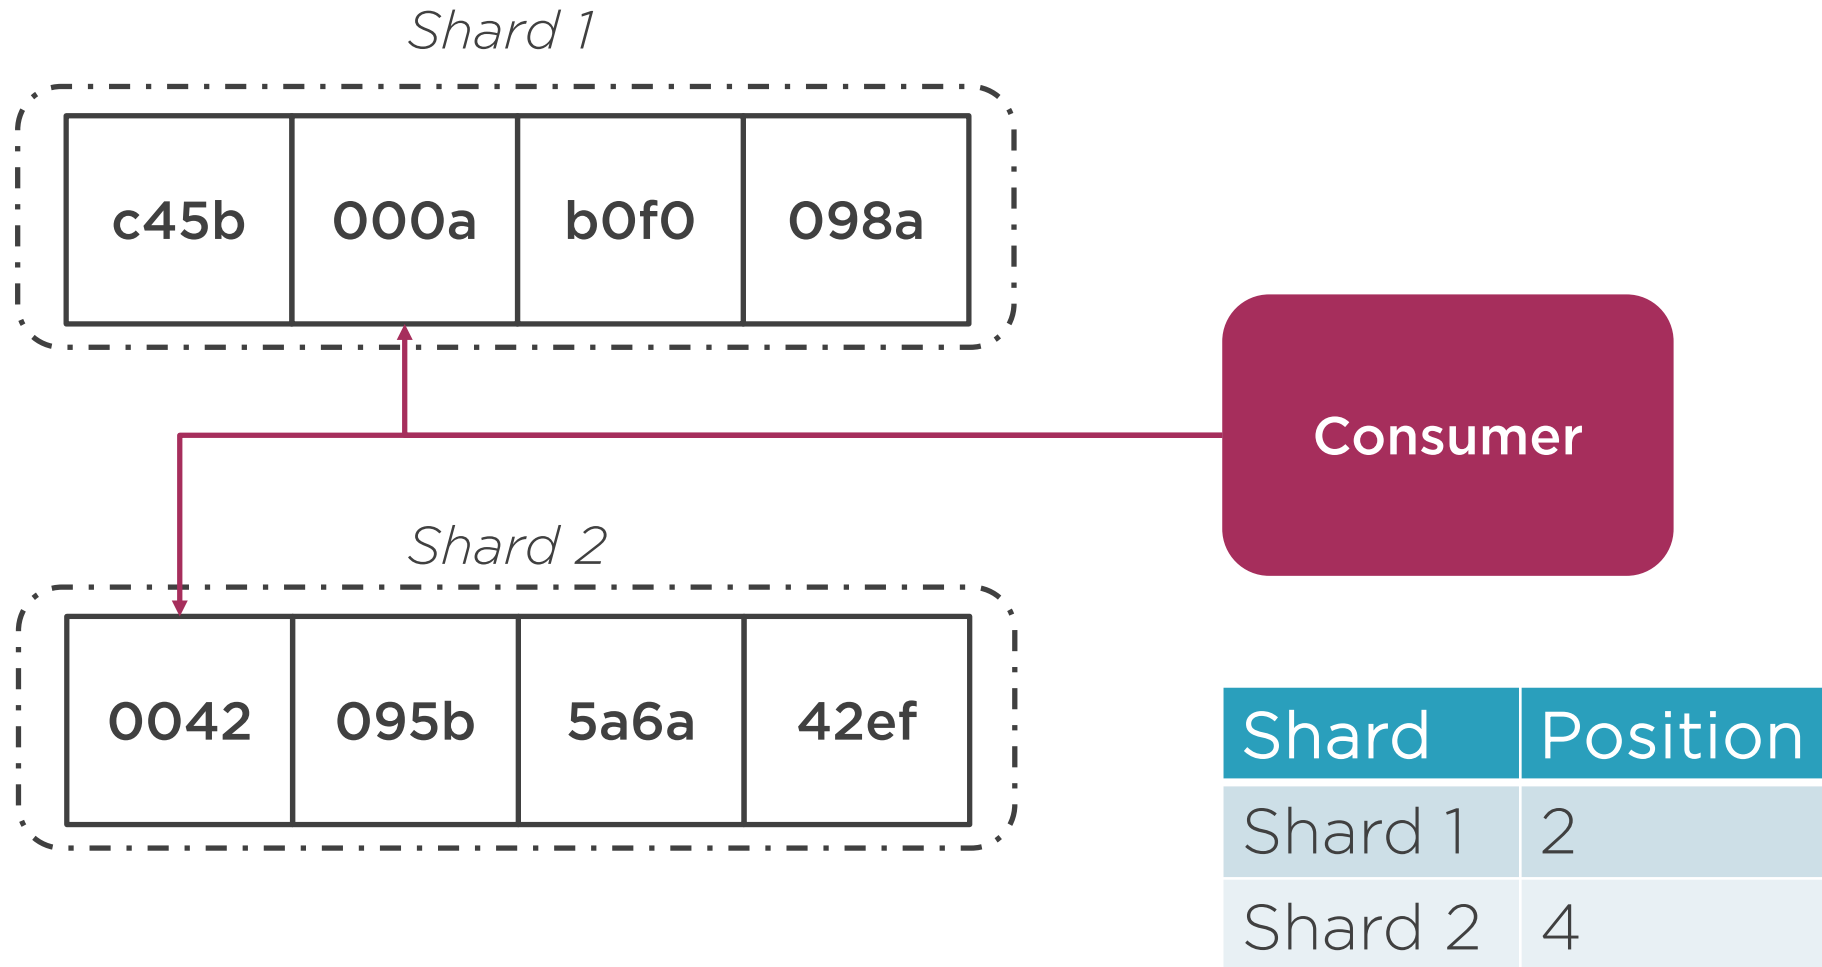
No need to implement this ourselves

- Other libraries: e.g. KCL

Reading Data from Kinesis



Reading Data from Kinesis

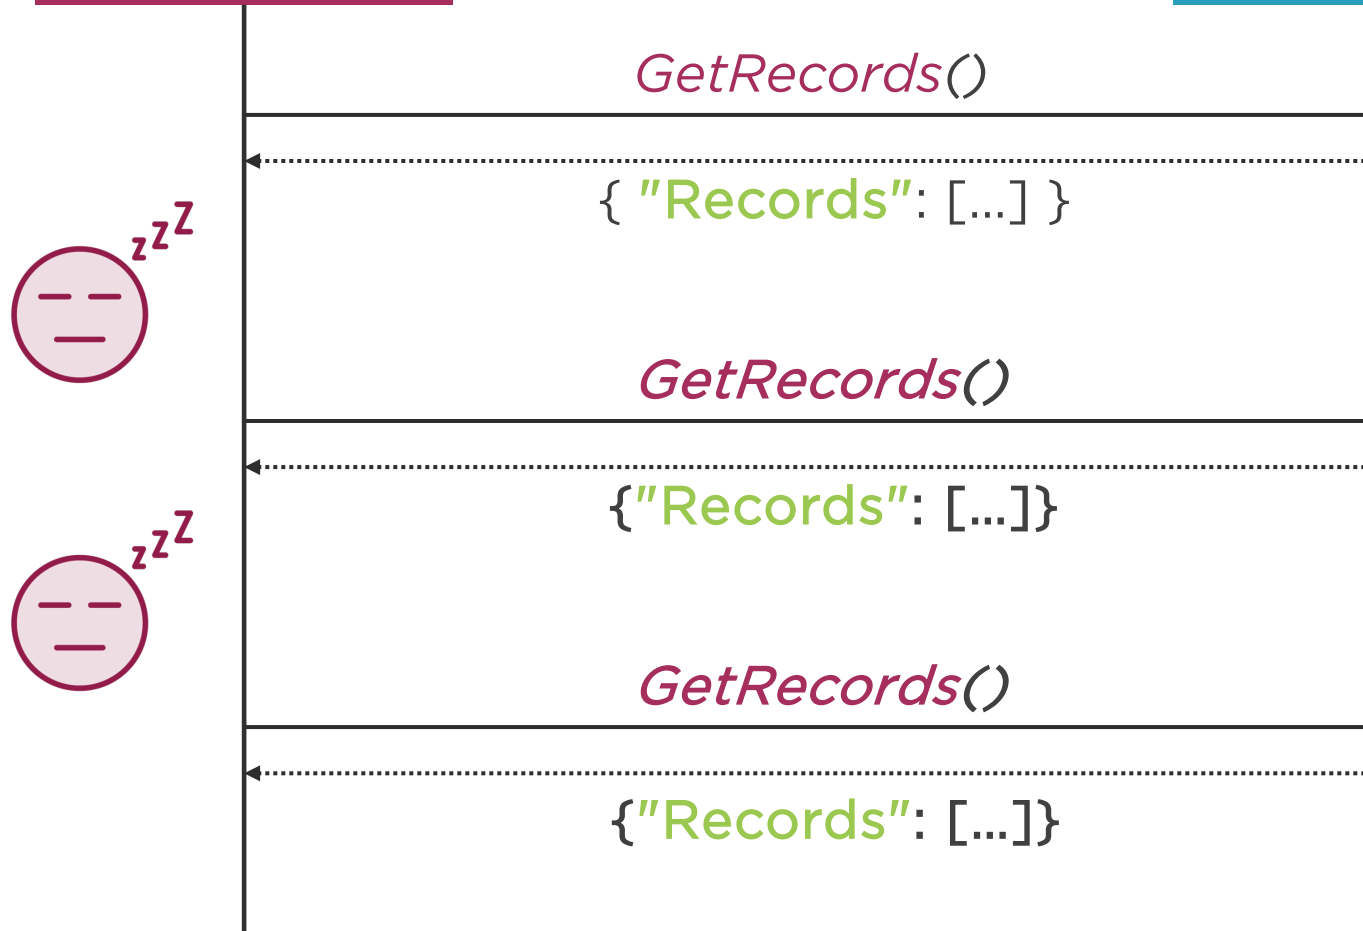


Reading from Kinesis



Consumer

Kinesis
shard



Shard Limitations



Amount of data


- One shard can return <2MB per second
- Divided among all consumers

Number of read requests


- 5 reads per shard per second
- Can call `GetRecords()` at most every 200 ms



Number of Read Requests



Number of Read Requests



Enhanced Fan-Out Consumer



Consumer

Kinesis
shard



SubscribeToShard()


{ "Records": [...] }

{ "Records": [...] }

{ "Records": [...] }



Enhanced Fan-Out Consumer



Amount of data

- Up to 2MB per second per consumer

Lower latency

- Around 70 ms

Number of consumer

- Up to 5 (not a hard limit)

Extra cost

- Per GB retrieved and per shard

SDK for Java 2.x



Fan-out consumer only for AWS SDK 2.x

- We were using SDK 1.x in this course
- Used in most existing apps

AWS SDK 2.x

- Major SDK rewrite
- Built for Java 8+

```
CompletableFuture<Record> future = getRecord();  
// Block until a result is ready  
future.get();  
// Do not block. Wait for result to be ready  
future.thenApply(r -> processRecord(r));
```

Java 8 CompletableFuture

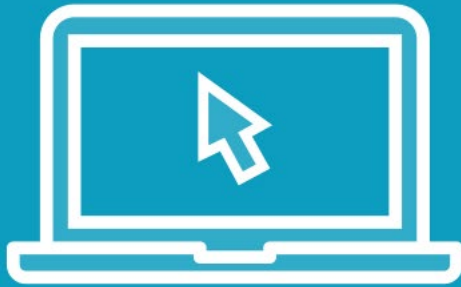
Class that represents a future result (e.g. HTTP request)

Can process results/errors asynchronously.

Can wait until a result is ready



Demo



Read tweets stream

Implement enhanced fan-out consumer



Using Kinesis Efficiently




A shard can process limited number of requests

- Will return throughput exceed exception
- **Solution:** use more shards

You can have more shards than necessary

- Will pay more than necessary
- **Solution:** reduce number of shards

Scaling Kinesis Stream



Scaling Kinesis API

SplitShard

Replace a single
shard
with two shards

MergeShards

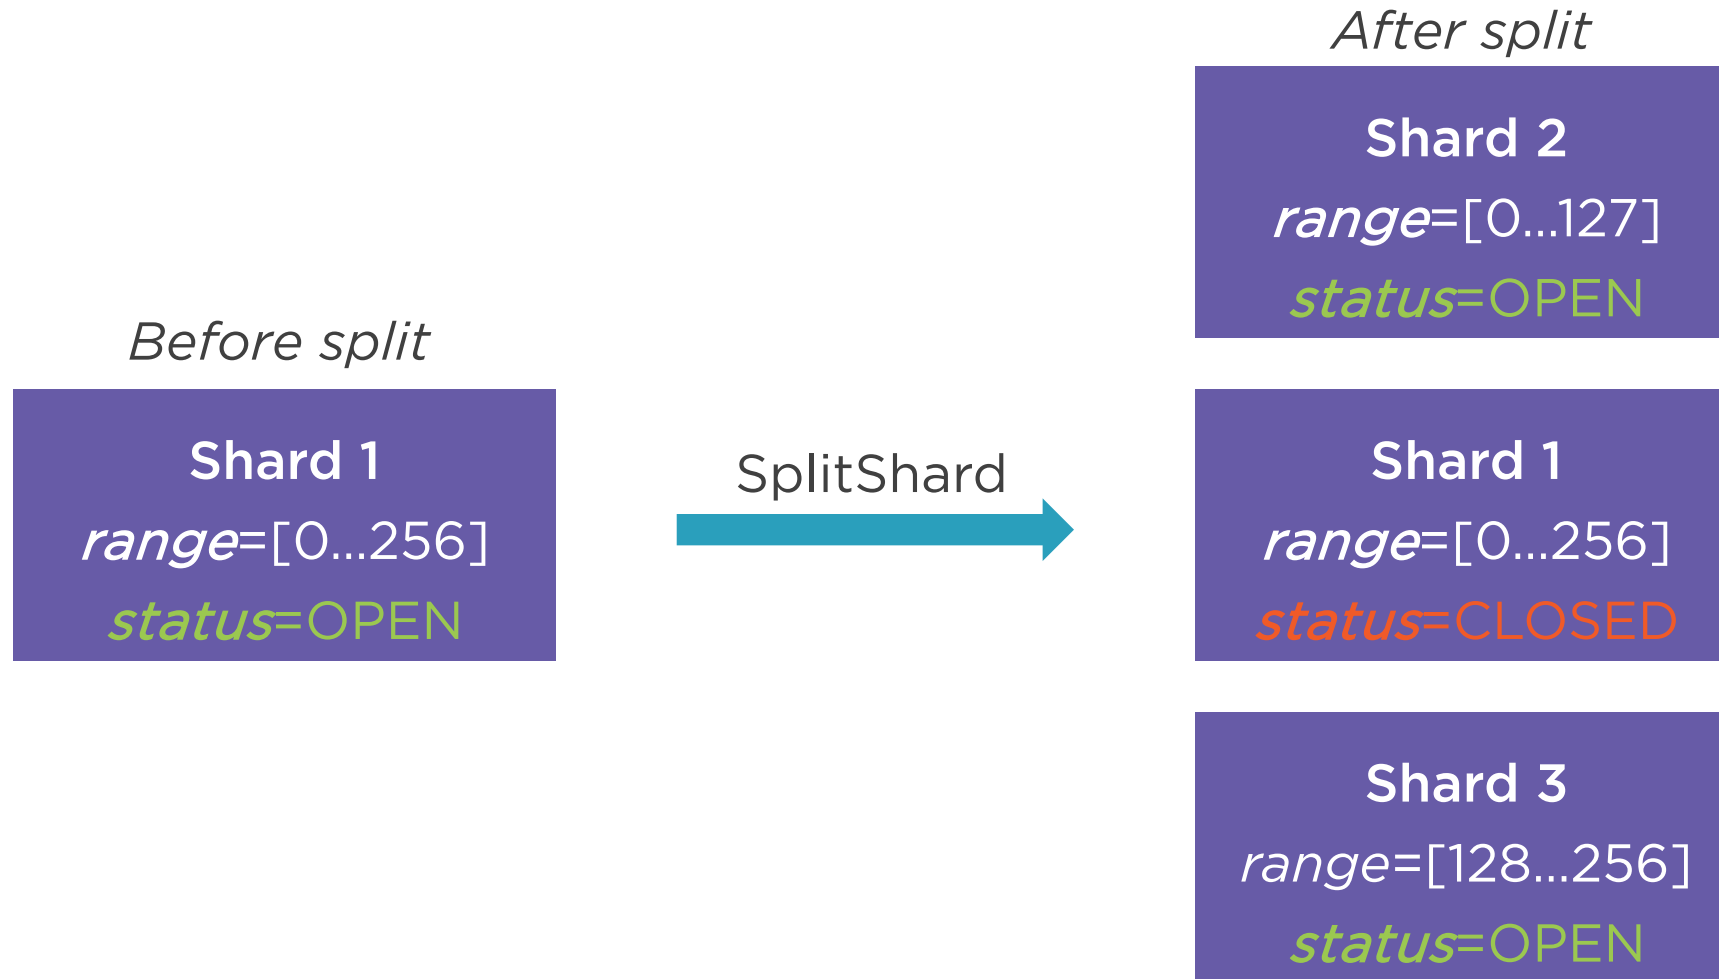
Replace two
shards with a
single shard

UpdateShardCount

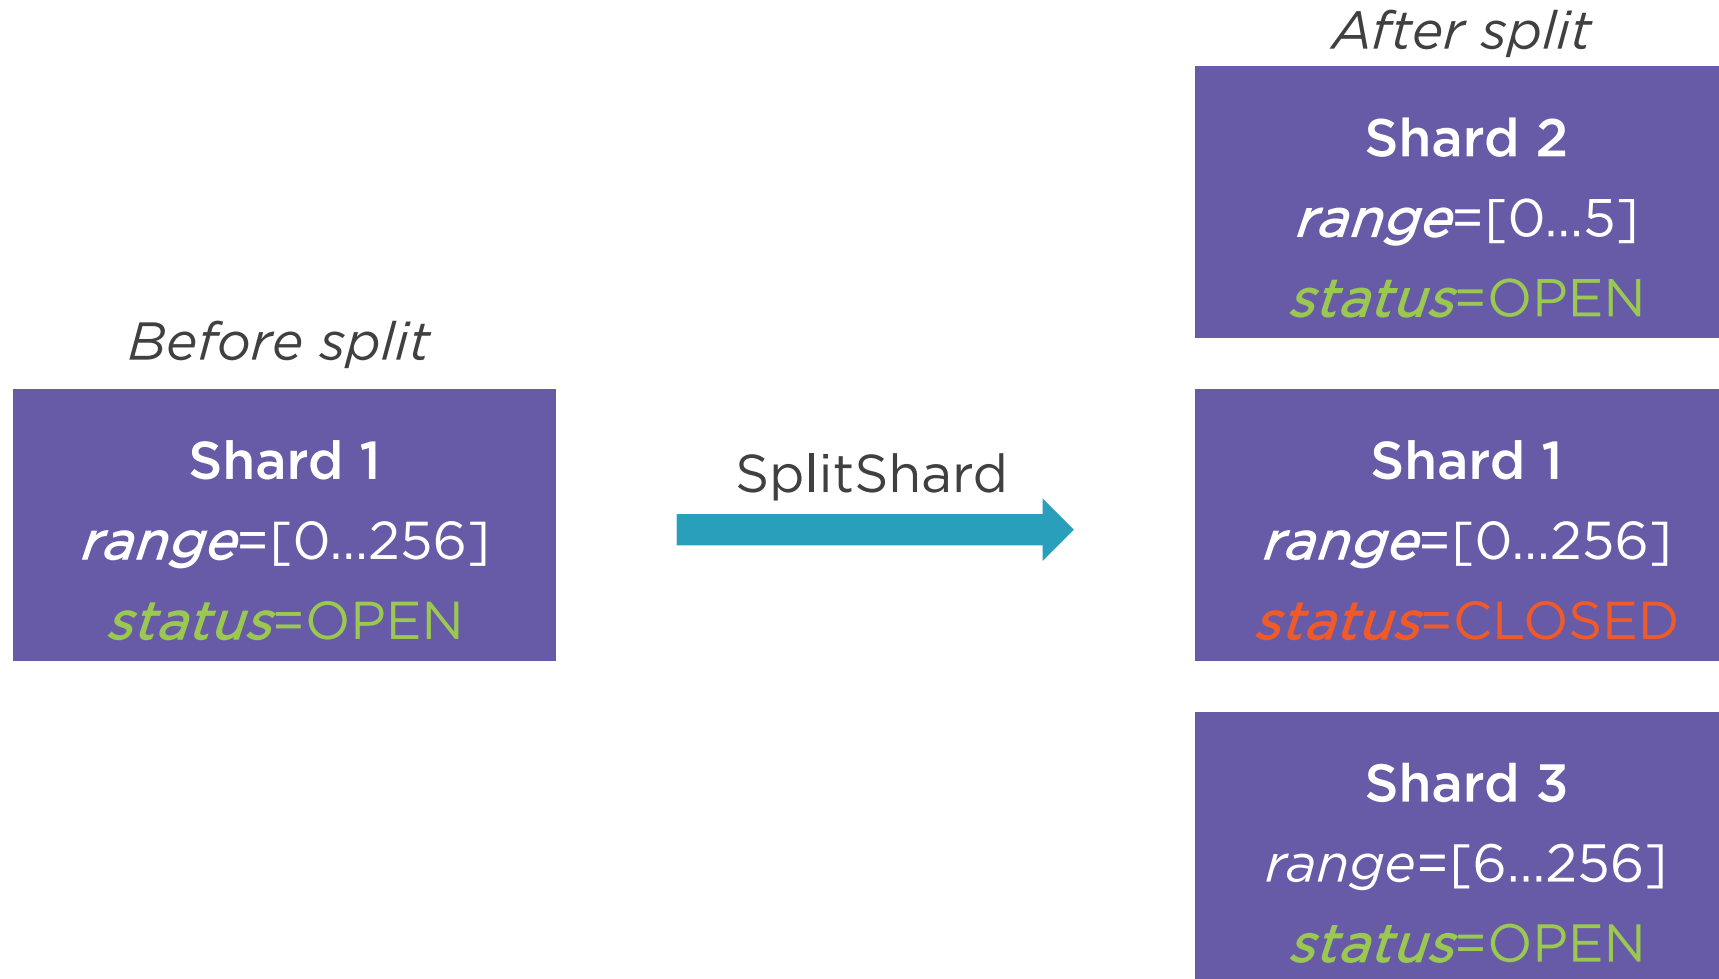
Scale up or down
to the specified
number of shards



Splitting a Shard



Splitting a Shard Unevenly



```
{  
  "NewStartingHashKey" : "127",  
  "ShardToSplit" : "shard-23",  
  "StreamName" : "tweetsStream"  
}
```

SplitShard API Call

Splits shard "shard-23" into two shards.



Merging Shards

Before merge

Shard 1
range=[0...127]
status=OPEN

Shard 2
range=[128...256]
status=OPEN

MergeShards



After merge

Shard 1
range=[0...127]
status=CLOSED

Shard 2
range=[128...256]
status=CLOSED

Shard 3
range=[0...256]
status=OPEN



```
{  
  "AdjacentShardToMerge" : "shard-2",  
  "ShardToMerge" : "shard-3",  
  "StreamName" : "tweetsStream"  
}
```

MergeShards API Call

Merges shard "shard-3" into "shard-2" and join their ranges

Returns HTTP 200

Shards must have adjacent ranges




```
{  
  "ScalingType": "UNIFORM_SCALING",  
  "StreamName": "tweetsStream",  
  "TargetShardCount": 100  
}
```

Setting Number of Shards

Change number of shards to 100.



Demo



Scale-up Kinesis stream

Scale-down Kinesis stream



Kinesis Limits



Data storage



Write per shard



Read from a shard



Data Storage Limitations



Number of shards:

- Up to 500 shards in Virginia, Oregon, Ireland
- Up to 200 shards elsewhere
- No a hard limit

Data retention period:

- 24 hours by default
- Can be up to 168 hours
- Hard limit

Data Writing Limitations




Shard throughput

- Up to 1MB
- **Solution:** Scale up


Number of transactions

- Up to 1,000 records per second
- **Solution:** Producer buffering


Write Request per Record



Write Request per Record



Write Request per Record



Read Limitations




Amount of data

- GetRecords returns <10MB
- One shard can return <2MB per second
- **Solution:** Scale up

Number of transactions

- 5 reads per shard per second
- **Solution:** Copy Kinesis data

Increasing Number of Transactions



Module Summary

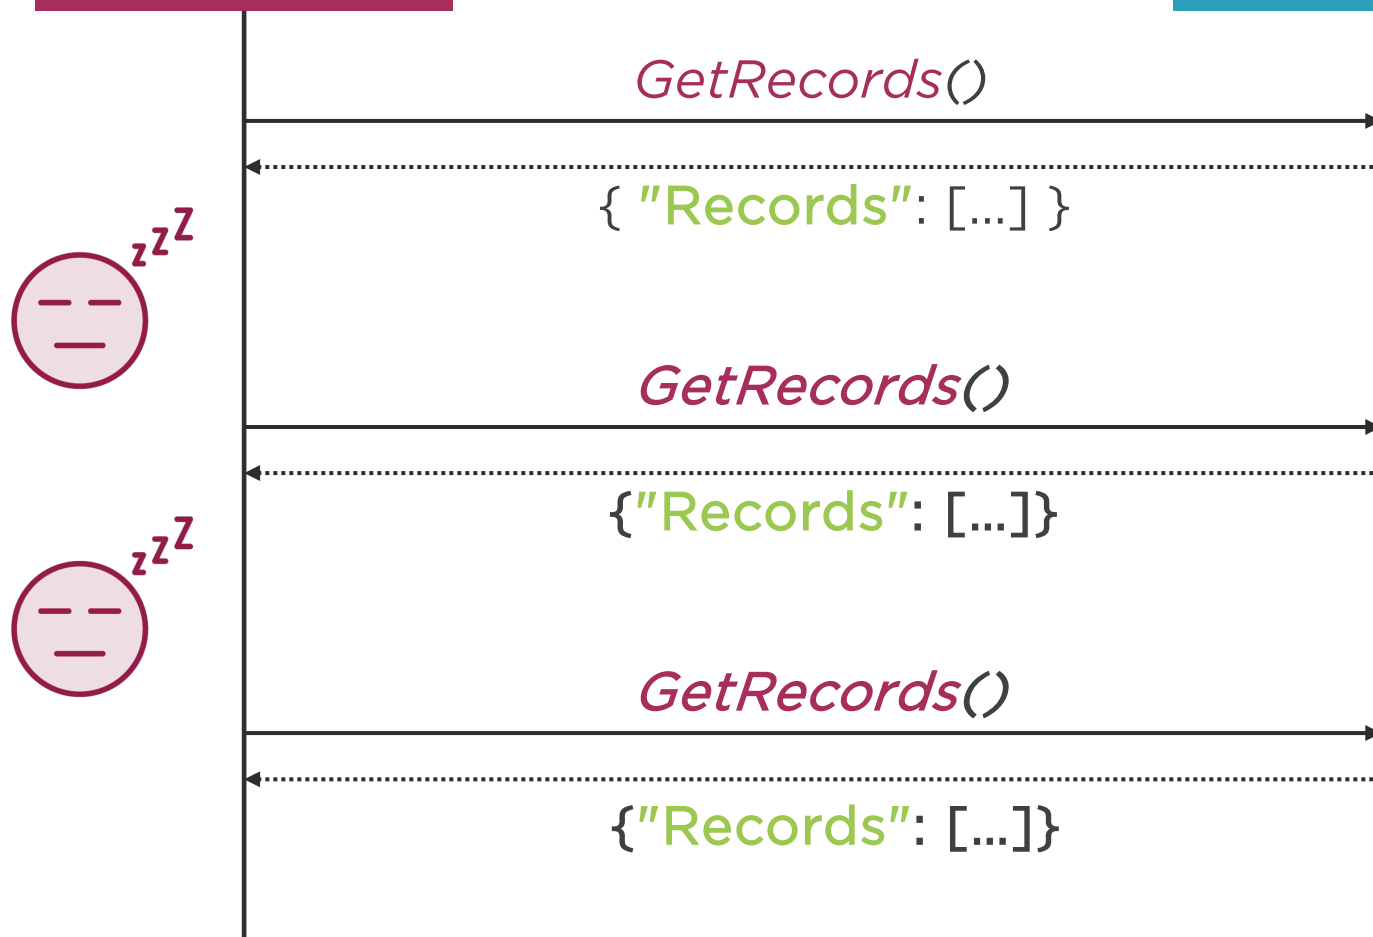


Module Summary



Consumer

Kinesis
shard



Module Summary



Consumer

Kinesis
shard



SubscribeToShard()

{ "Records": [...] }

{ "Records": [...] }

{ "Records": [...] }



Module Summary Cont.

