

Developing Applications Using Kinesis Client Library



Ivan Mushketyk

@mushketyk brewing.codes



What Are We Going to Implement

Tweets

mushketyk 13:45

The weather is nice!

anonymous 13:46

@mushketyk Stop
procrastinating!

anonymous 13:46


@mushketyk Go record a
course already!




language	time	count
en	13:45	1
en	13:46	2



How Are We Going to Implement It



Kinesis Producer Library



Why KPL?

Simple API

Better
performance

Lower cost

Errors handling

Monitoring



Writing Data to Kinesis

```
KinesisProducerConfiguration config = new  
KinesisProducerConfiguration();  
  
KinesisProducer = new KinesisProducer(config);  
  
kinesisProducer  
    .addUserRecord(  
        "stream",  
        partitionKey,  
        recordBytes);
```



Writing Data to Kinesis

```
KinesisProducerConfiguration config = new  
KinesisProducerConfiguration();
```

```
KinesisProducer = new KinesisProducer(config)
```

```
ListenableFuture<UserRecordResult> f = kinesisEnabledProducer  
    .addUserRecord(  
        "stream",  
        partitionKey,  
        recordBytes);
```

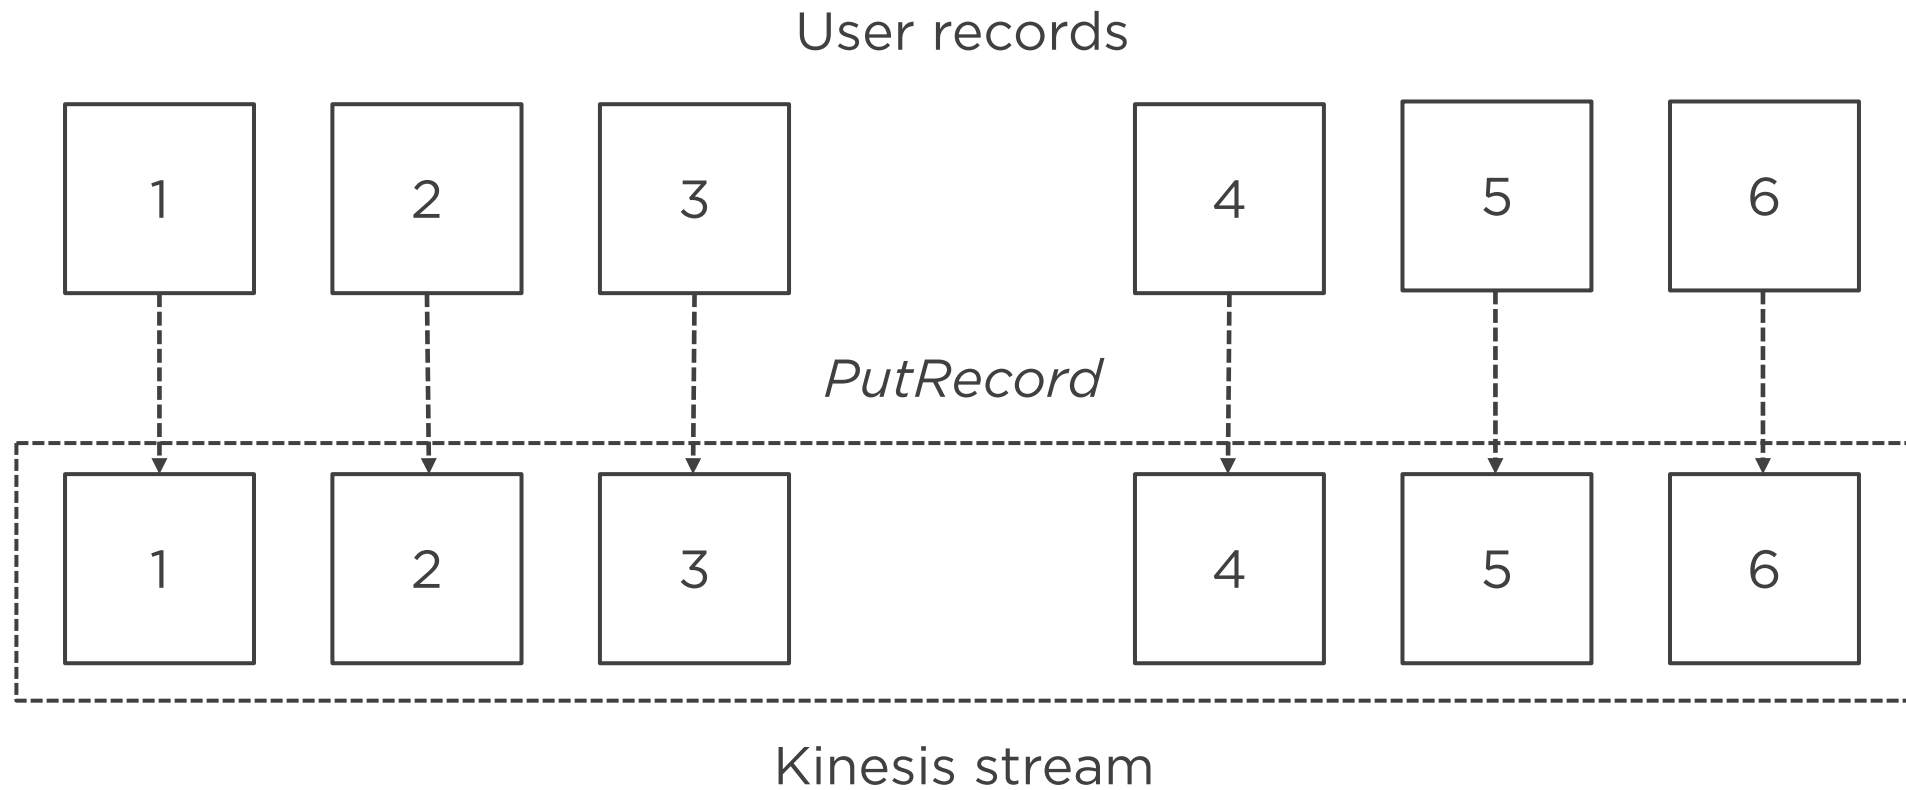


Process Future

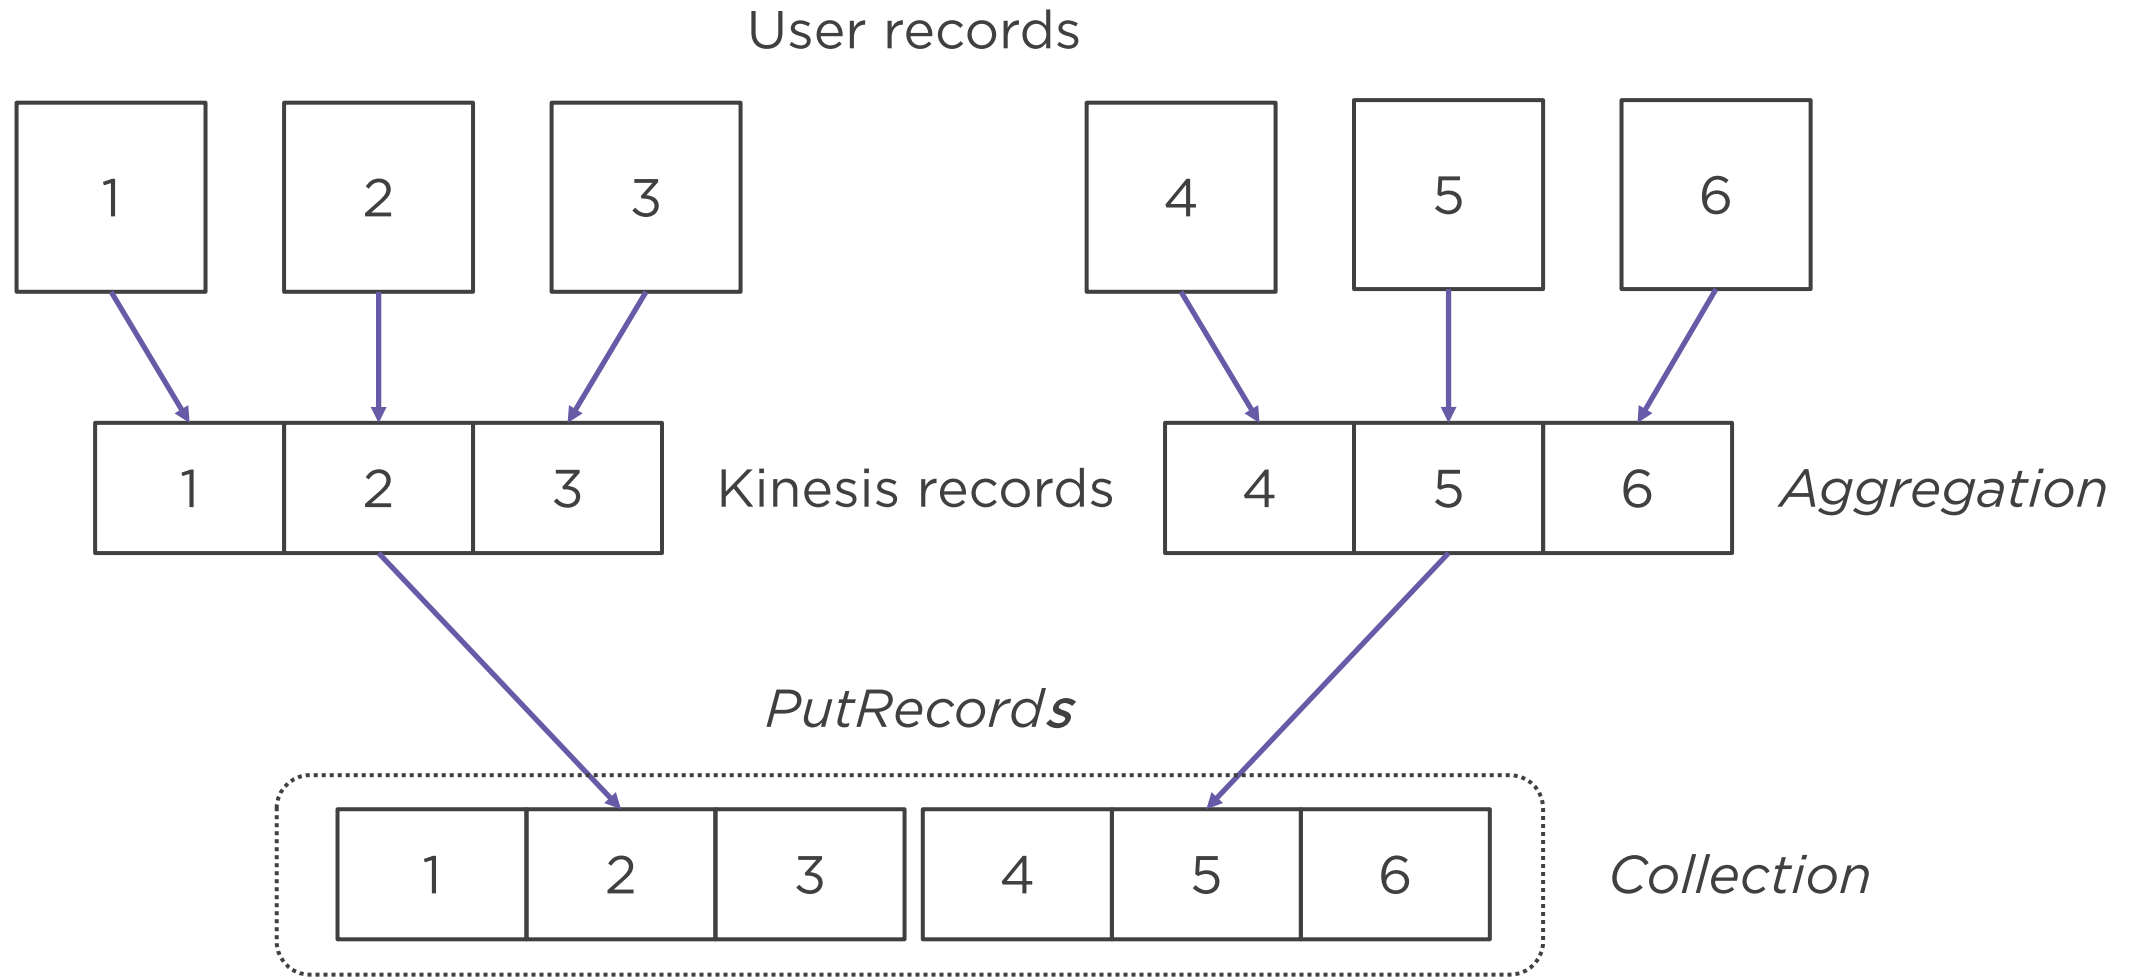
```
ListenableFuture<UserRecordResult> f = ...  
  
// Synchronously  
UserRecordResult result = f.get();  
  
// Asynchronously  
Futures.addCallback(f, new FutureCallback<UserRecordResult>() {  
    public void onFailure(Throwable t) {...}  
    public void onSuccess(UserRecordResult result) {...}  
});
```



Storing Records



Storing Records



Records Aggregation



Up to 1000 records per shard per second

1M PutRecord calls costs \$0.014

Can be tricky

- Consumer should have access to same records
- Limit latency

Can be disabled

```
config
```

```
.setRecordMaxBufferedTime(3000)
```

```
.setMaxConnections(1)
```

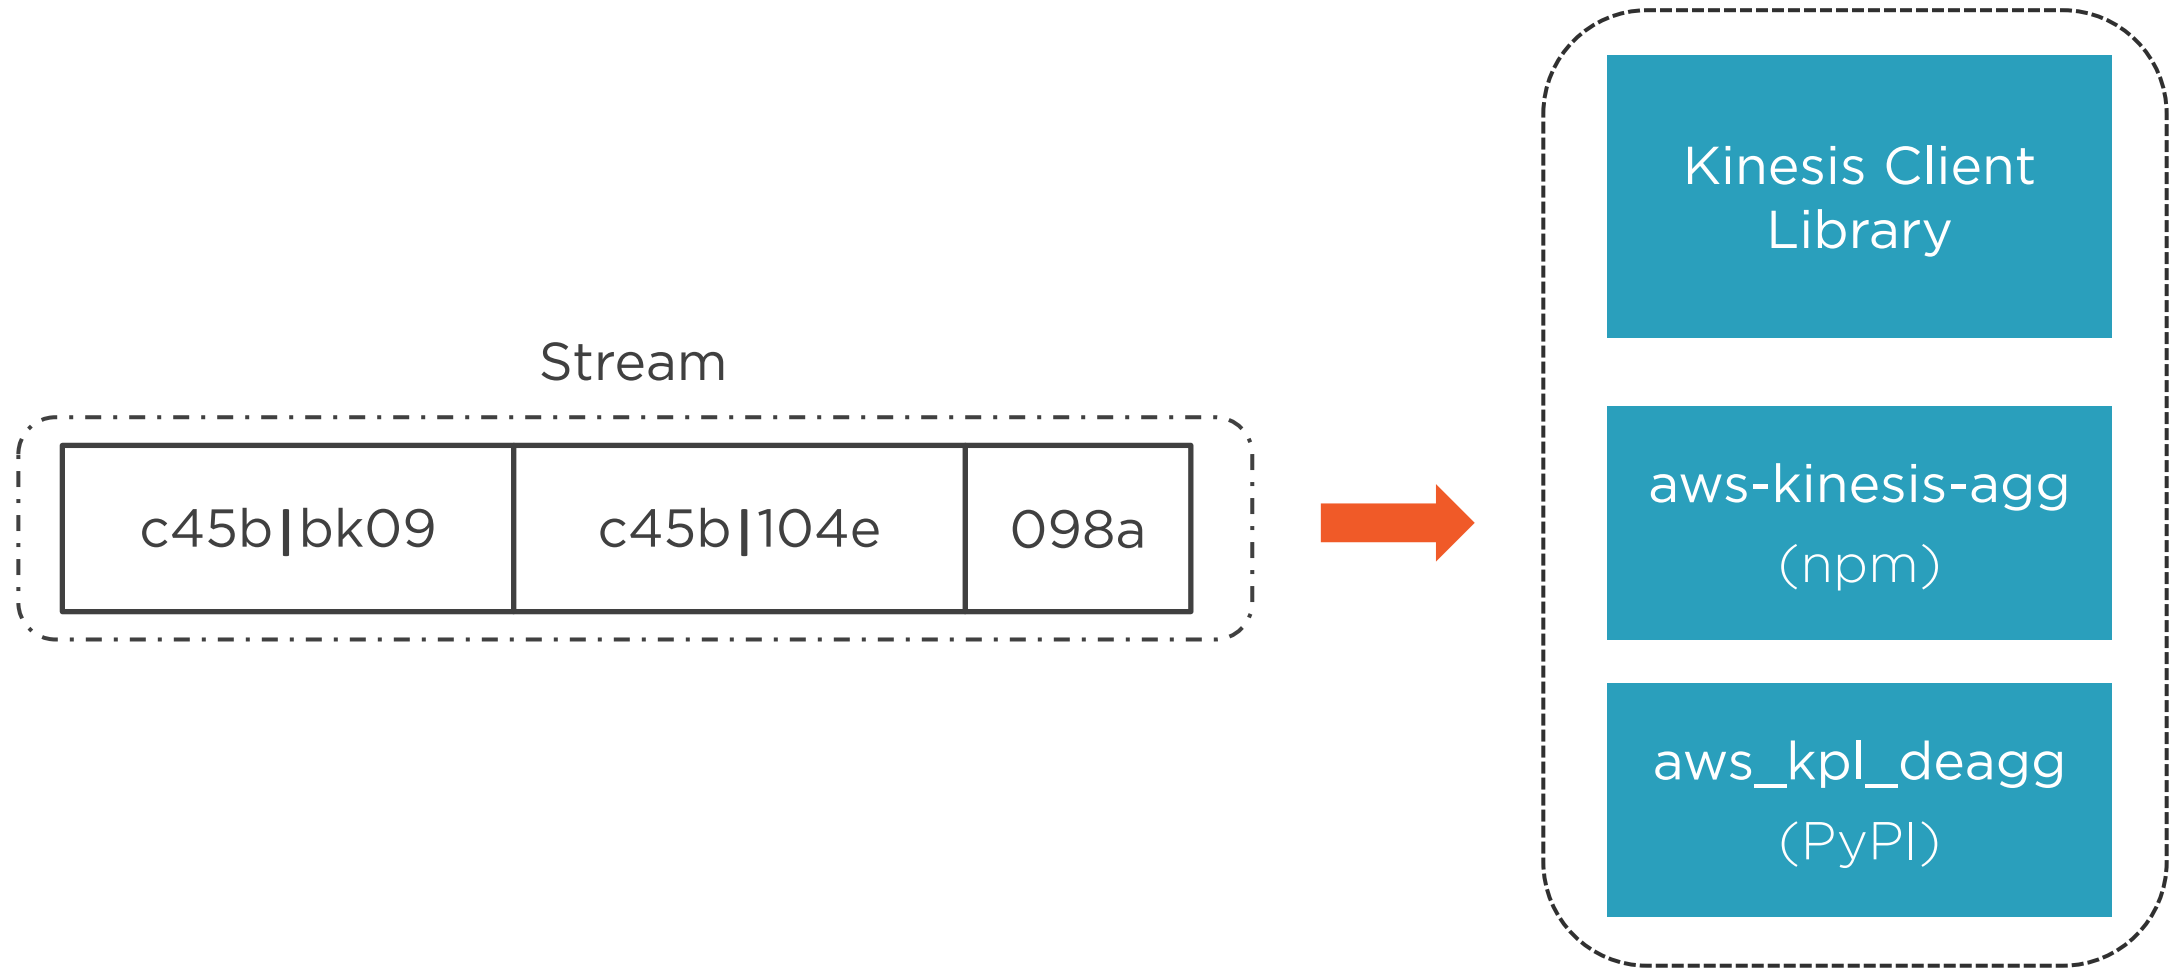
```
.setRequestTimeout(60000)
```

```
.setRegion("us-west-1");
```

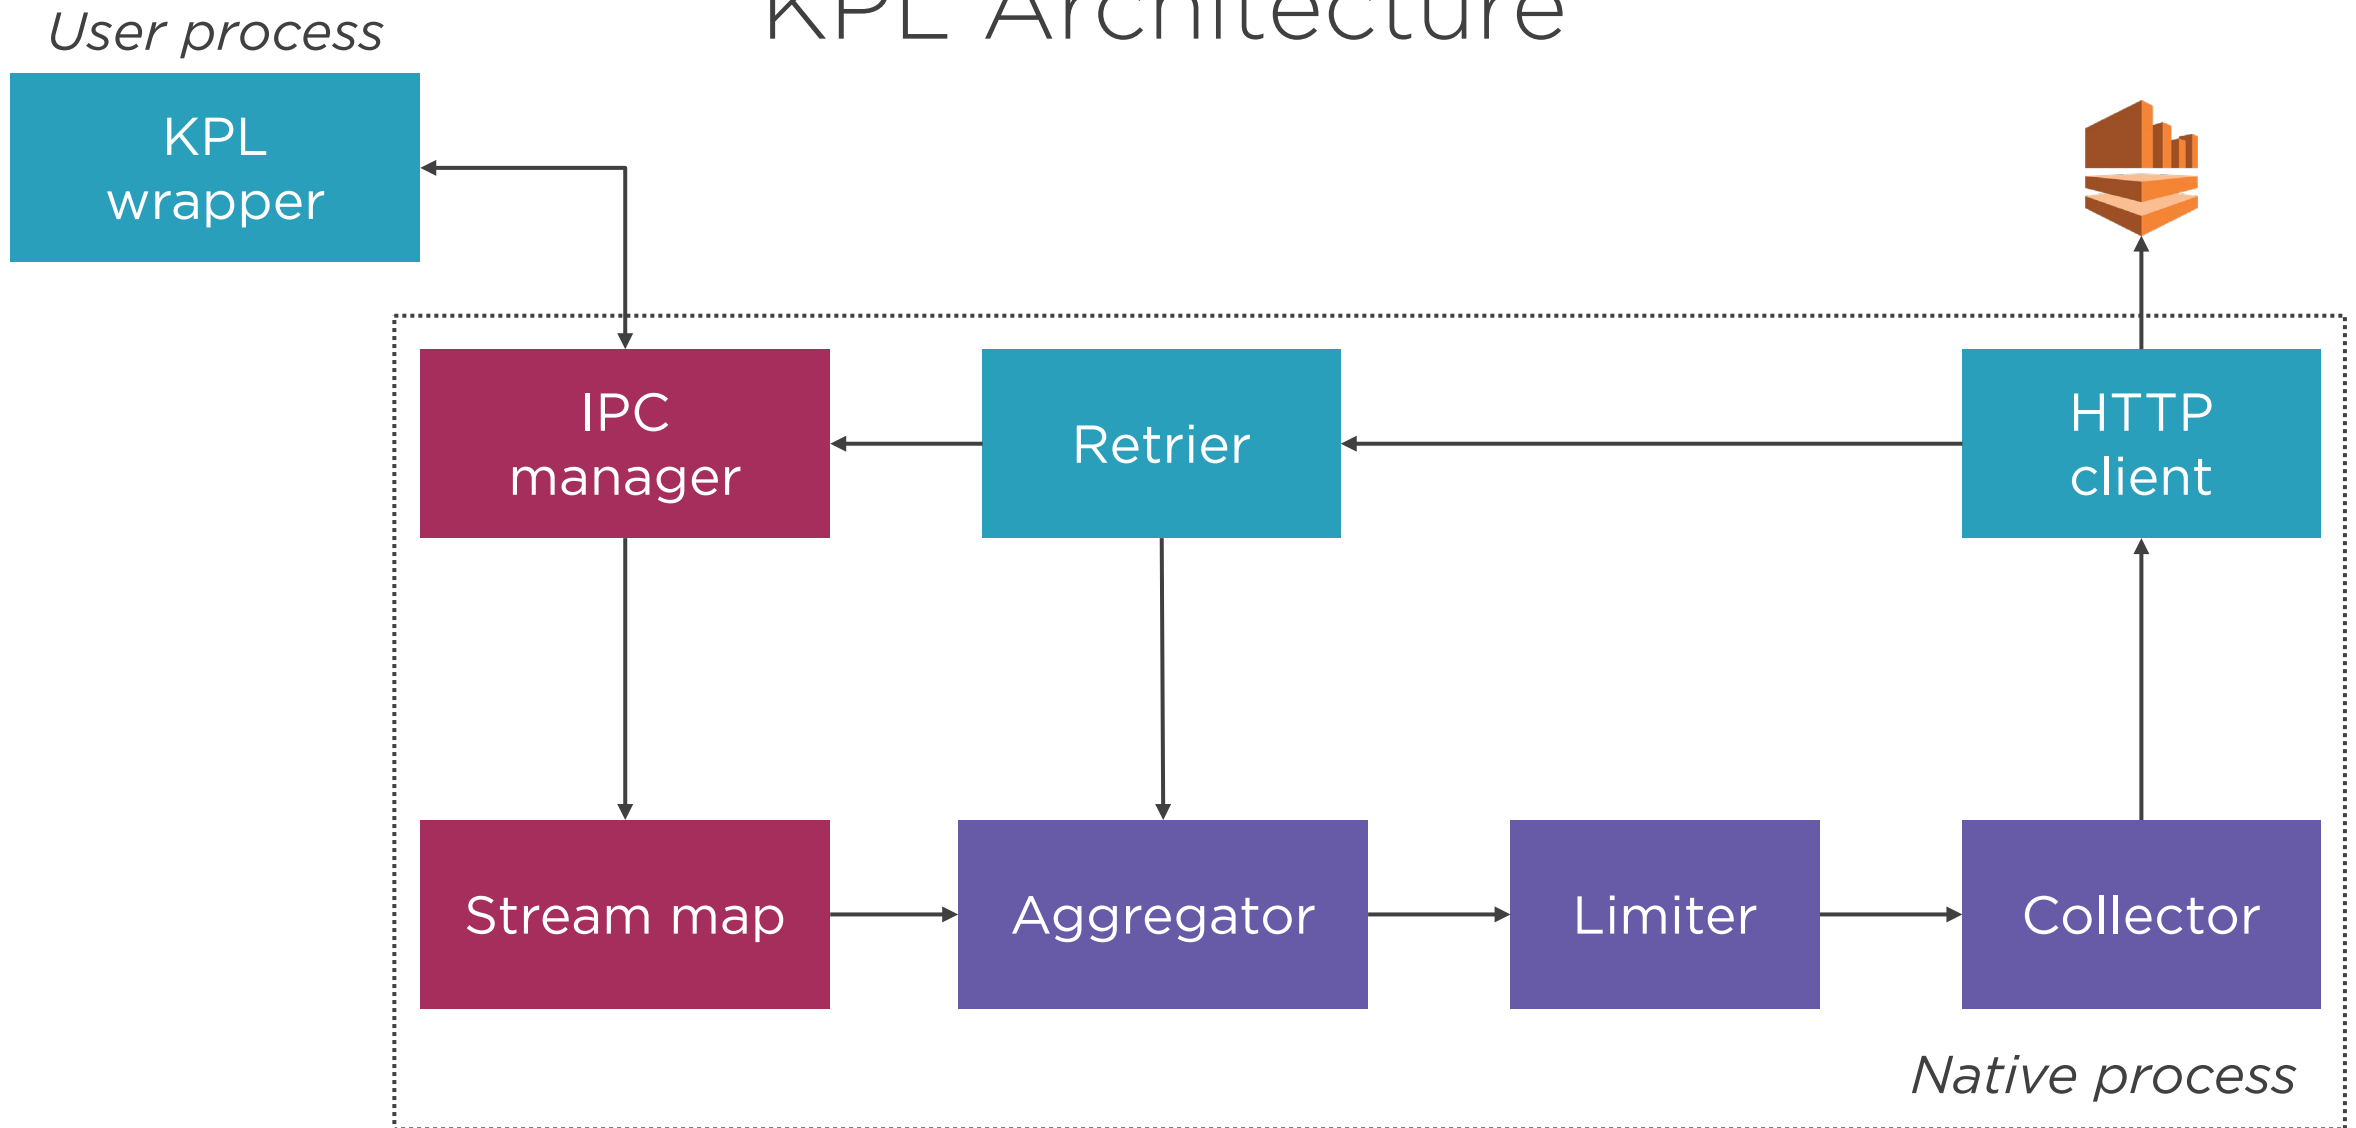
- ◀ How long an user record should be buffered, higher value results in more records being aggregated
- ◀ Degree of parallelism for HTTP requests
- ◀ Timeout for a single HTTP request. Higher timeout for slow connections
- ◀ AWS region



Decoding KPL Records




KPL Architecture



Aggregating with Multiple Shards

Aggregation



Shard 1
range=[0...127]

Aggregation



Shard 2
range=[128...256]

```
PutRecord {  
    "ExplicitHashKey": "string"  
}
```



Detailed Metrics

**User Record
Received**

**User Records
Pending**

**Kinesis Records
Put**


Buffering Time

**User Records per
Kinesis Record**

Errors by Code



Using Kinesis Client Library



Why KCL?

Simple API

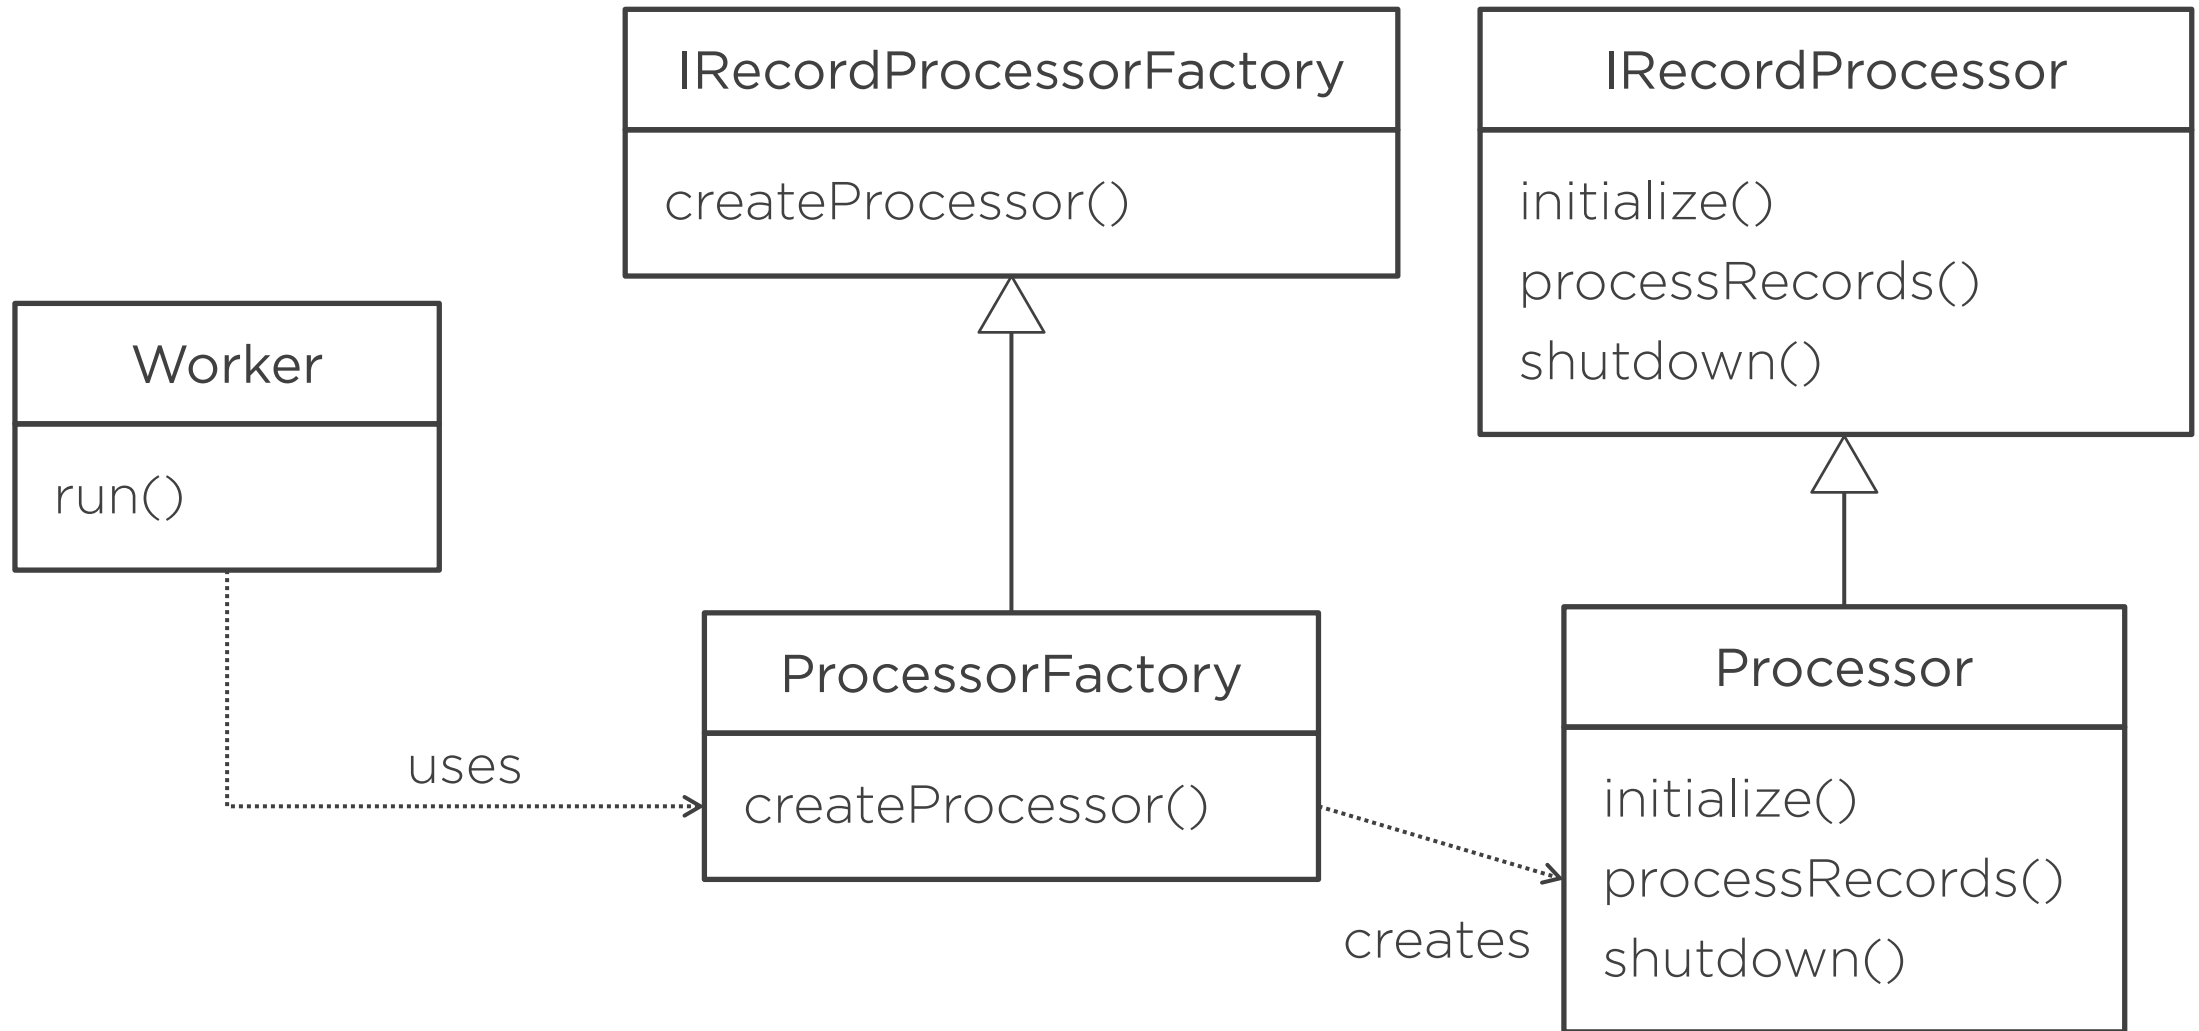
Handles resharding

Errors handling

Monitoring



Kinesis Client Library



Using Worker

```
KinesisClientLibConfiguration config = new
KinesisClientLibConfiguration(
    "tweets-processor", // Application name
    "tweets-stream",    // Stream name
    new DefaultAWSCredentialsProviderChain(),
    "worker-1"          // Name of the worker instance
);

IRecordProcessorFactory recordProcessorFactory = ...

Worker worker = new Worker.Builder()
    .config(config)
    .recordProcessorFactory(recordProcessorFactory)
    .build();
```



Configuring Worker

```
KinesisClientLibConfiguration config = ...  
  
// Starting position in a stream  
// Can be: LATEST, AT_TIMESTAMP, TRIM_HORIZON  
config.withInitialPositionInStream(  
    InitialPositionInStream.LATEST);  
  
// How often to query for new data  
config.withIdleTimeBetweenReadsInMillis(200);
```




Processing Records


```
public void processRecords(ProcessRecordsInput input) {  
    for (Record record : processRecordsInput.getRecords()) {  
        String partKey = record.getPartitionKey();  
        ByteBuffer data = record.getData();  
        String seq = record.getSequenceNumber();  
  
        if (record instanceof UserRecord)  
            String subSeq = record.getSubSequenceNumber();  
    }  
}
```




Scaling with KCL



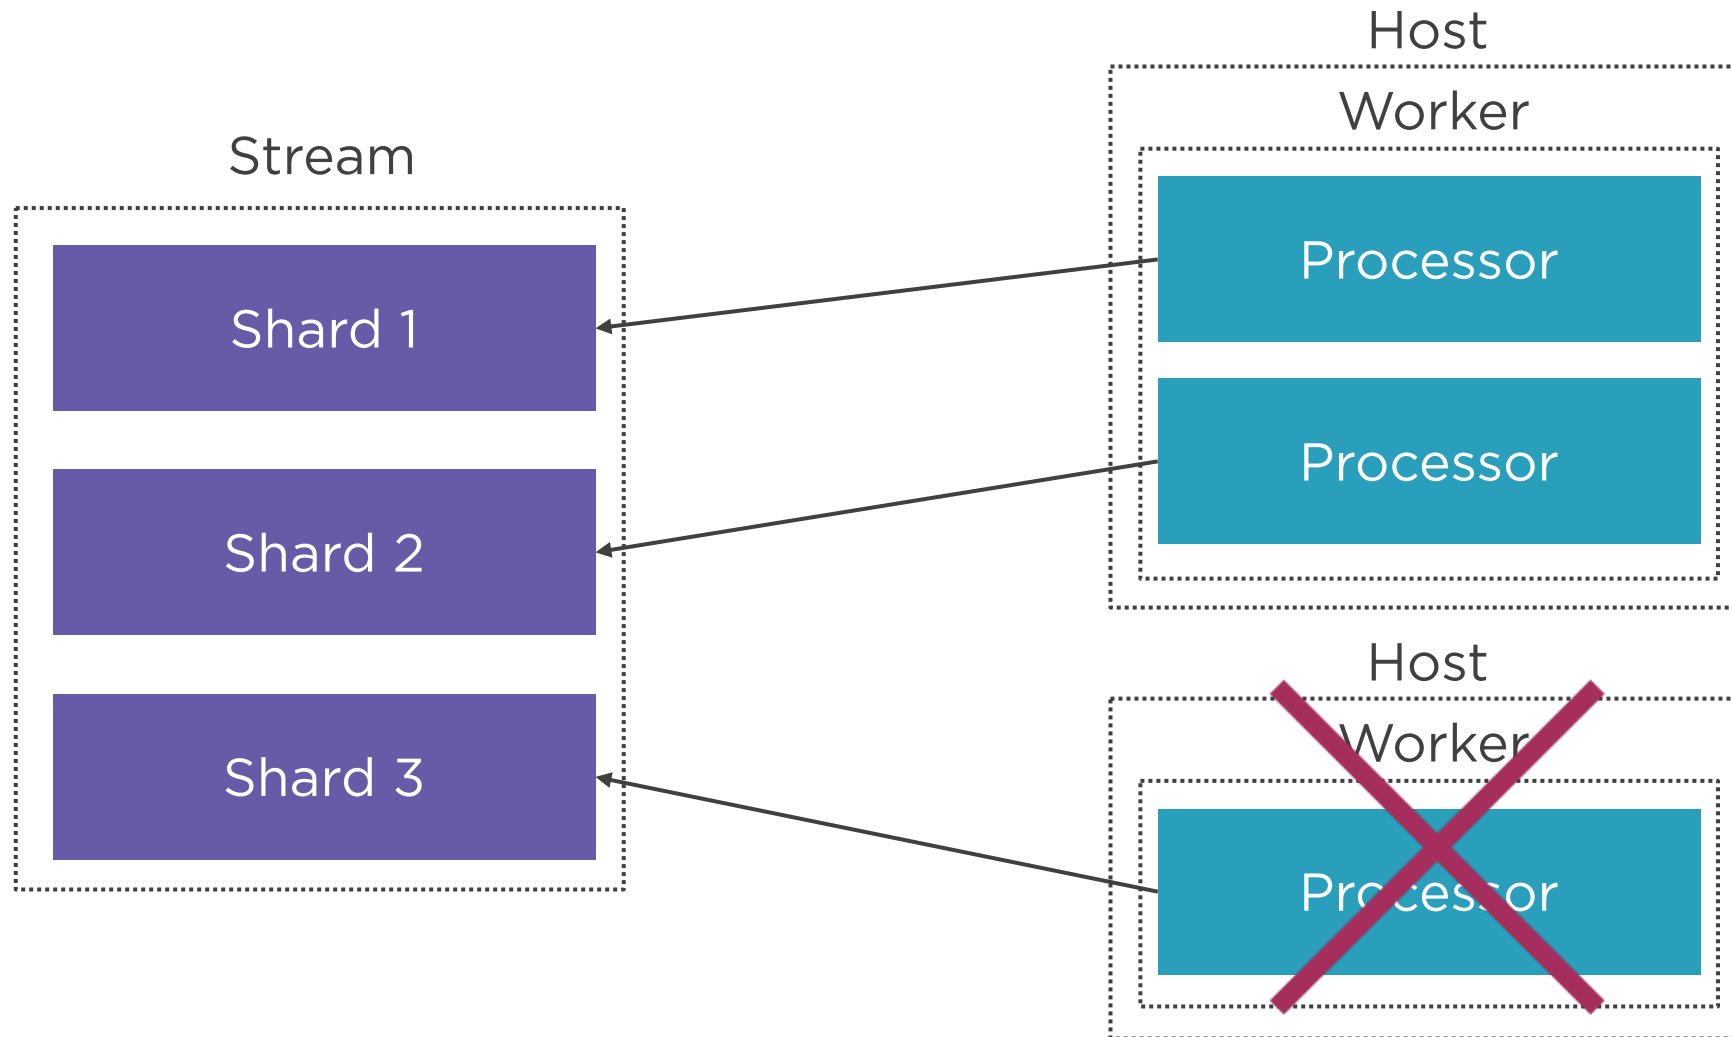
Scaling with KCL




Scaling with KCL



Scaling with KCL



Scaling with KCL



Kinesis Client Library does not launch additional hosts. It only starts new processors.



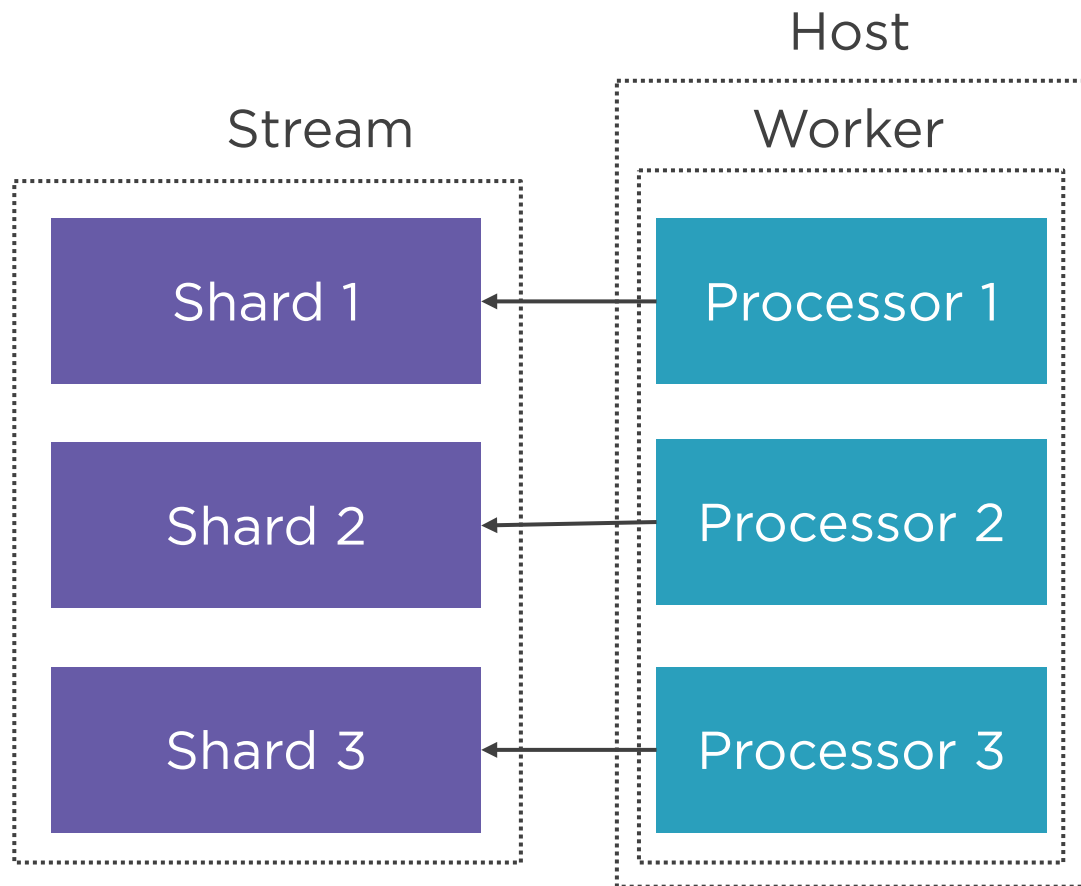
Demo



Implement simple application with KCL
Read tweets written with KPL



Fault Tolerance with KCL



shard	seq	subSeq	owner
1	100	1	processor_1
2	200	0	processor_2
3	134	2	processor_3

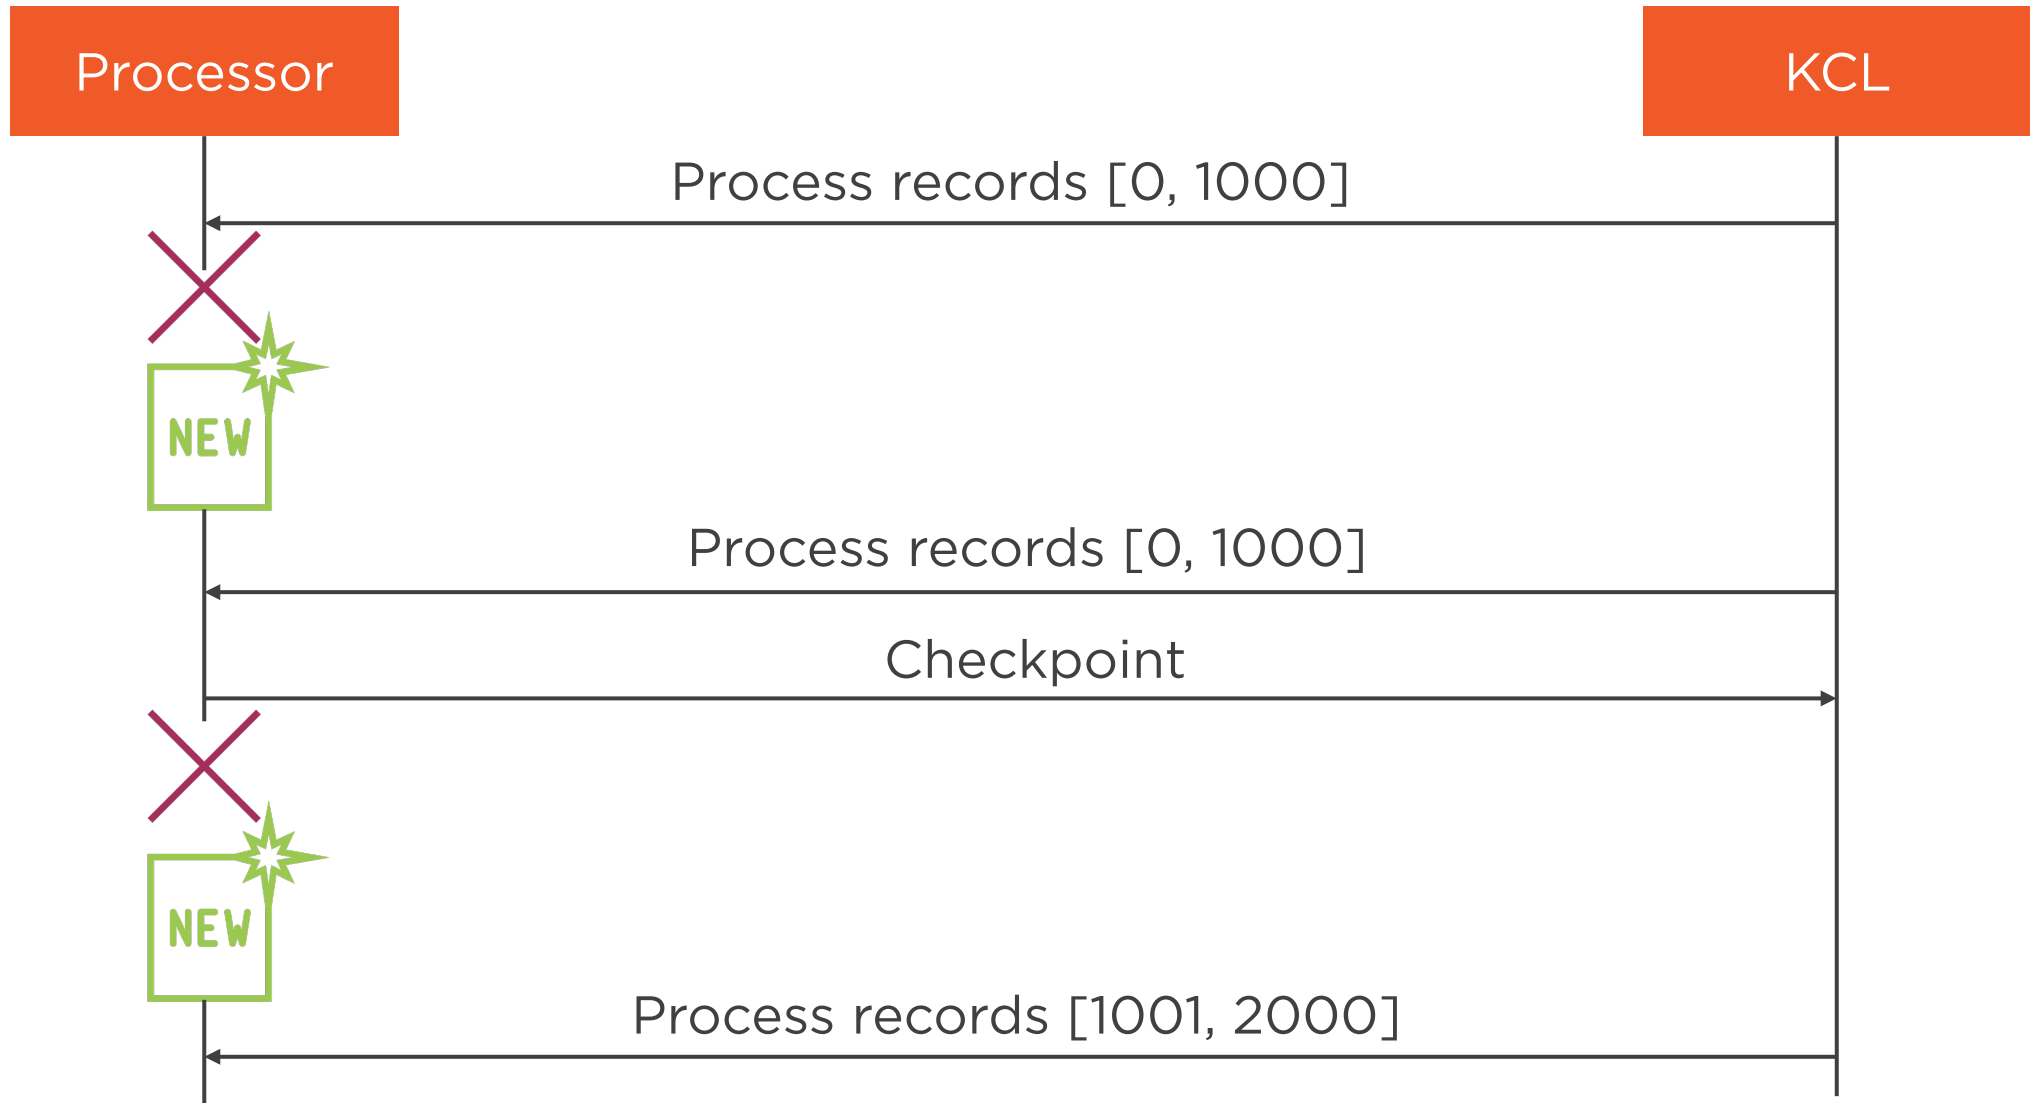


Checkpointing

```
public void processRecords(ProcessRecordsInput input) {  
    IRecordProcessorCheckpoint checkpoint =  
        input.getCheckpoint();  
}  
  
// Checkpoint last record  
checkpoint.checkpoint();  
// Checkpoint specified record  
checkpoint.checkpoint(record);  
// Checkpoint sequence number  
checkpoint.checkpoint(sequenceNumber);  
// Checkpoint aggregated record from KPL  
checkpoint.checkpoint(sequenceNumber, subSequenceNumber);
```



KCL Checkpointing



Demo




Add fault tolerance

Use checkpointing correctly



Tweet-counting Application



Counting the Number of Tweets

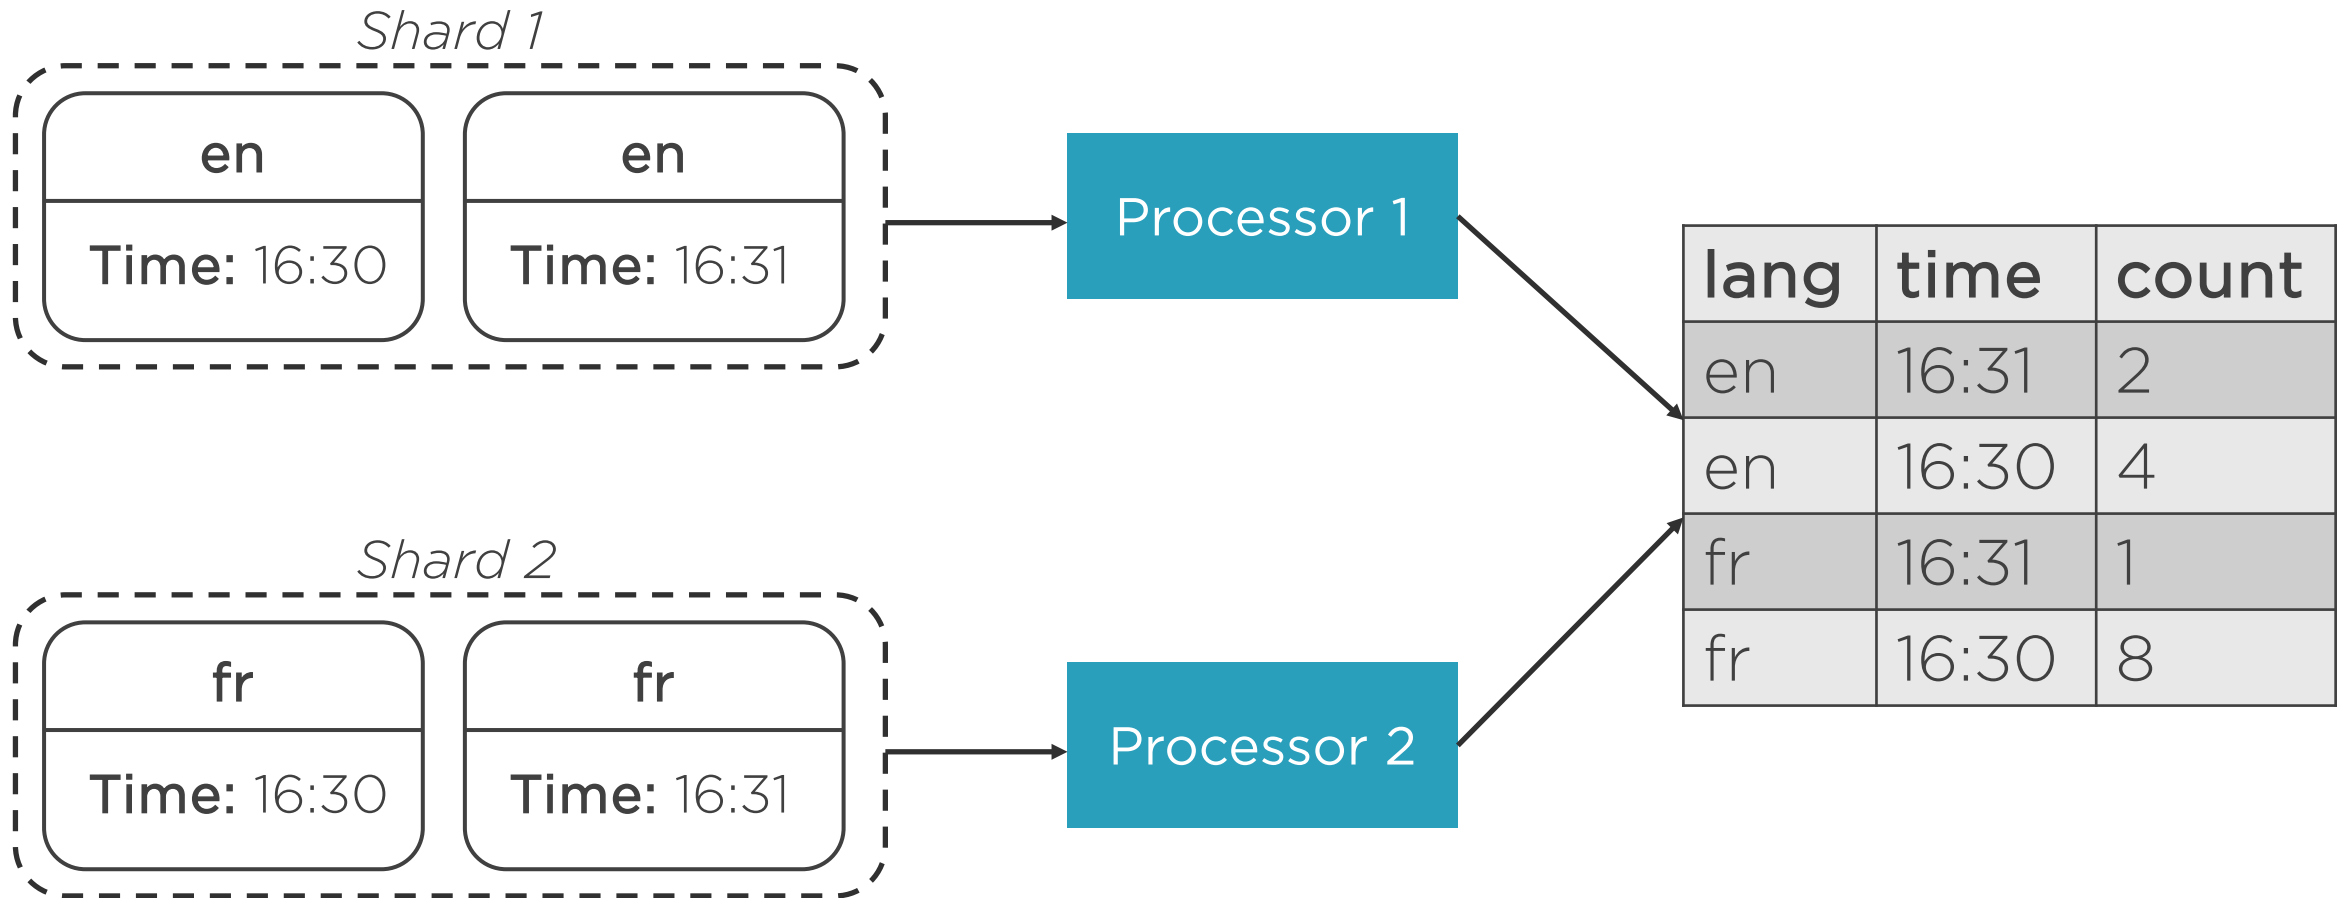


lang	time	count
en	16:30	2

lang	time	count
en	16:31	2




Processing Multiple Shards



Event Time



Processing time
System time




Event time
Time when an event was created

Idempotence

The property of certain operations that can be applied multiple times without changing the result



Idempotent Log Processing



Prepare Checkpointing

```
void initialize(InitializationInput input) {  
    input.getPendingCheckpointSequenceNumber();  
}  
// Prepare checkpoint  
checkpointer.prepareCheckpoint();  
checkpointer.prepareCheckpoint(record);  
checkpointer.prepareCheckpoint(sequenceNumber);  
checkpointer.prepareCheckpoint(  
    sequenceNumber,  
    subSequenceNumber);
```




Demo



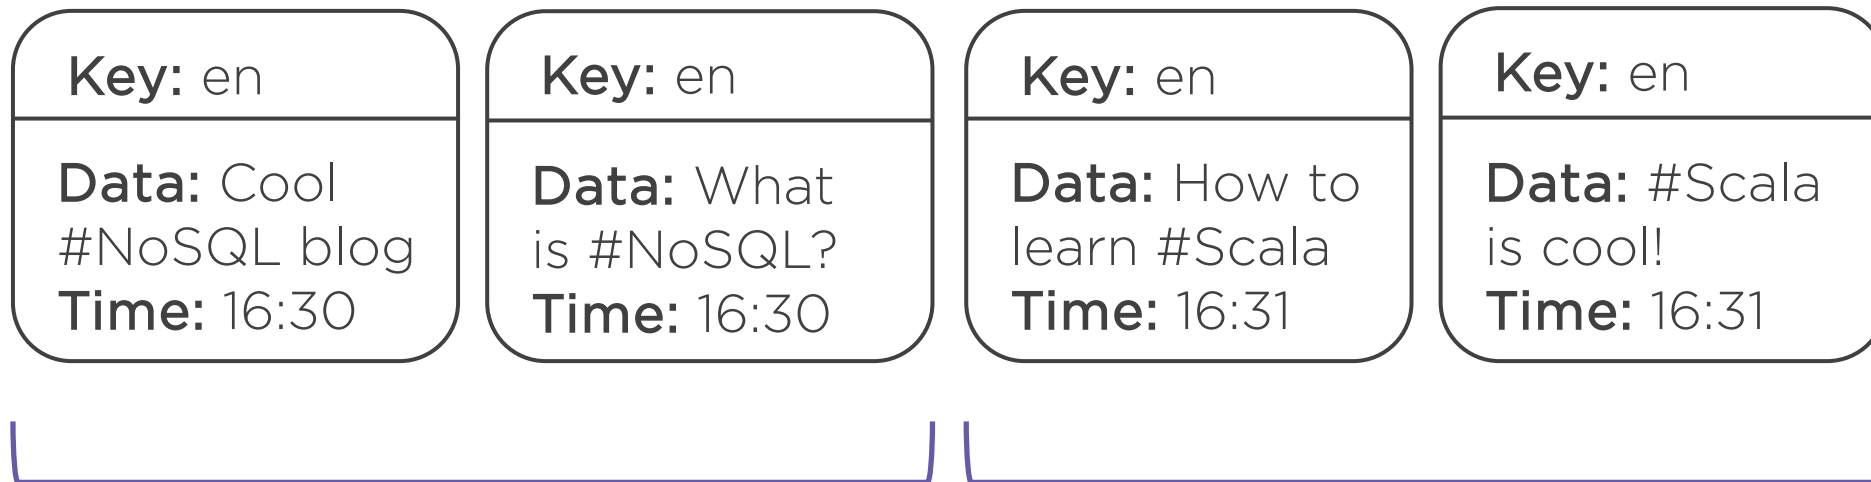
Implement tweet-counting application



Checkpointing with Resharding



Recording Shard Name



lang	time	shard	count
en	16:30	0	2

lang	time	shard	count
en	16:31	0	2



Demo



Fix tweet-counting application



Is There a Better Way?



What are some problems with our approach?

Why KCL is not enough?

What Is Missing – Too Much Code



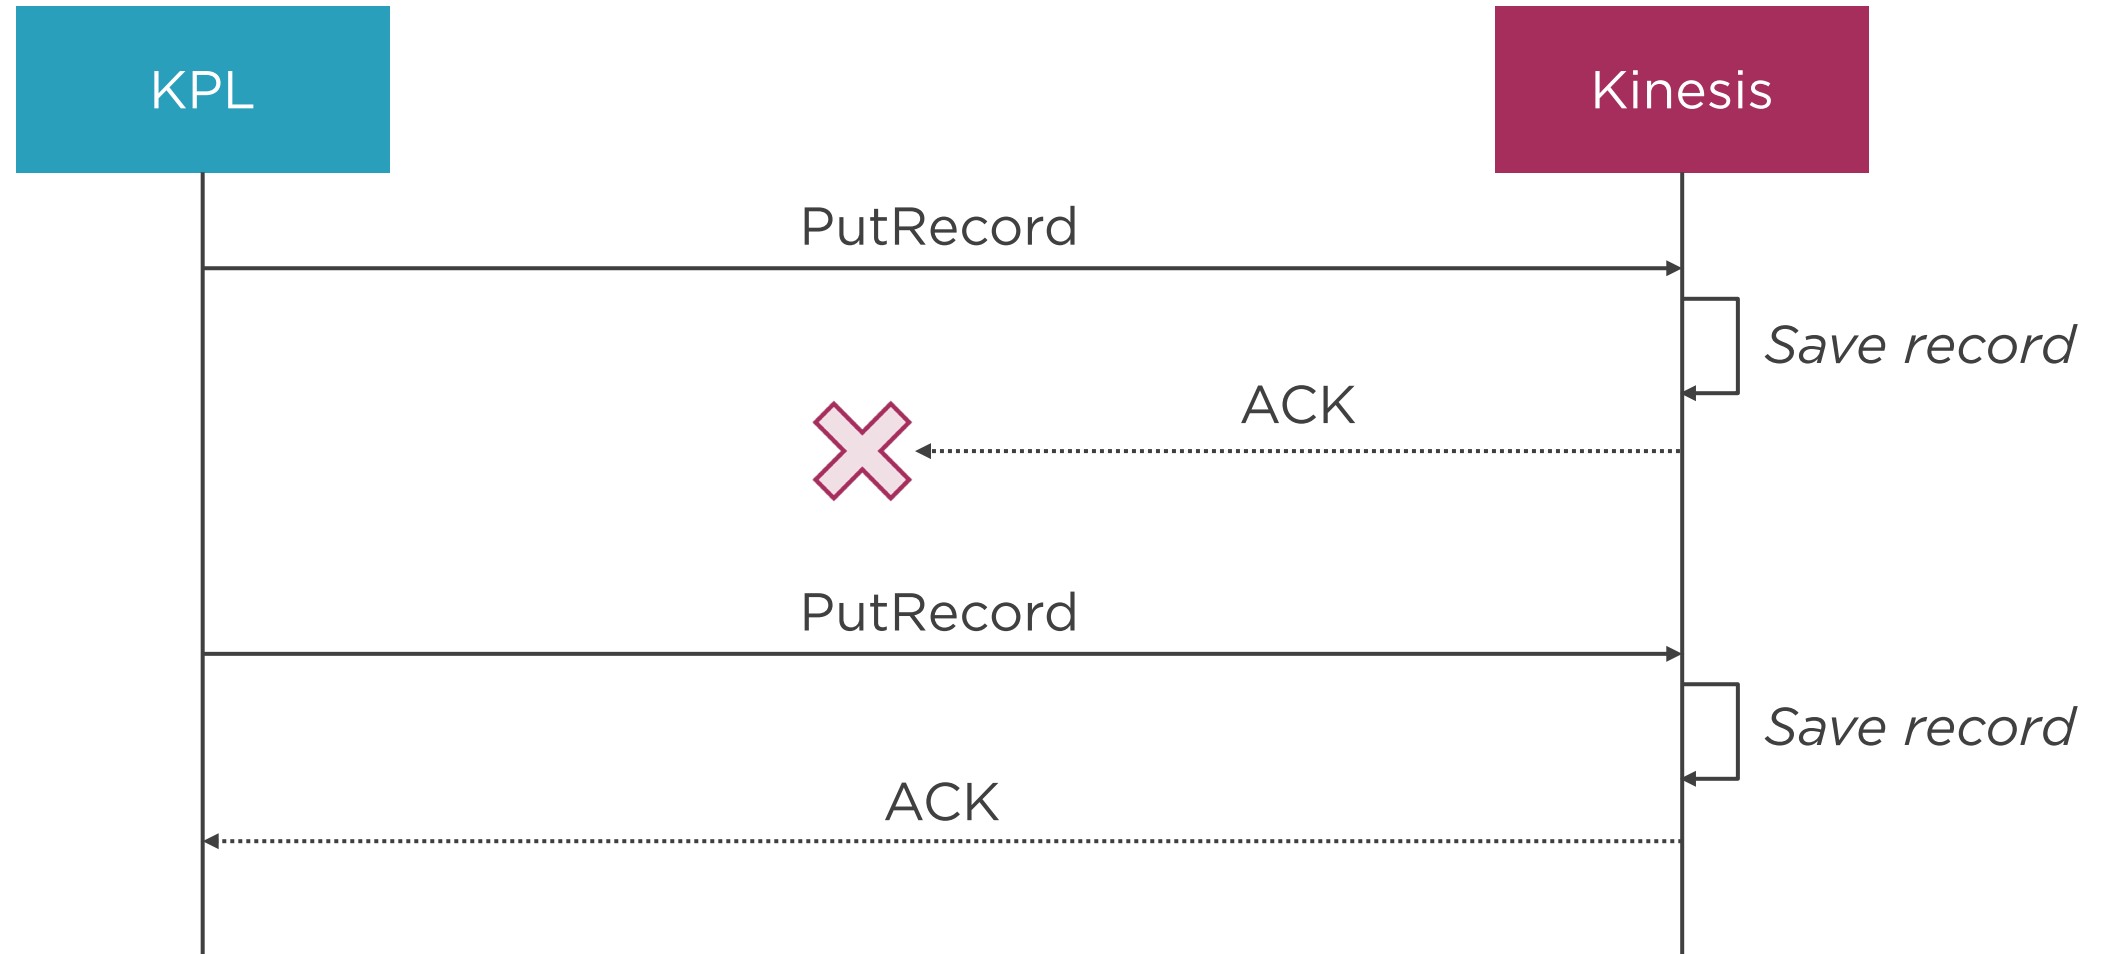
Implemented simple example

Still quite low-level

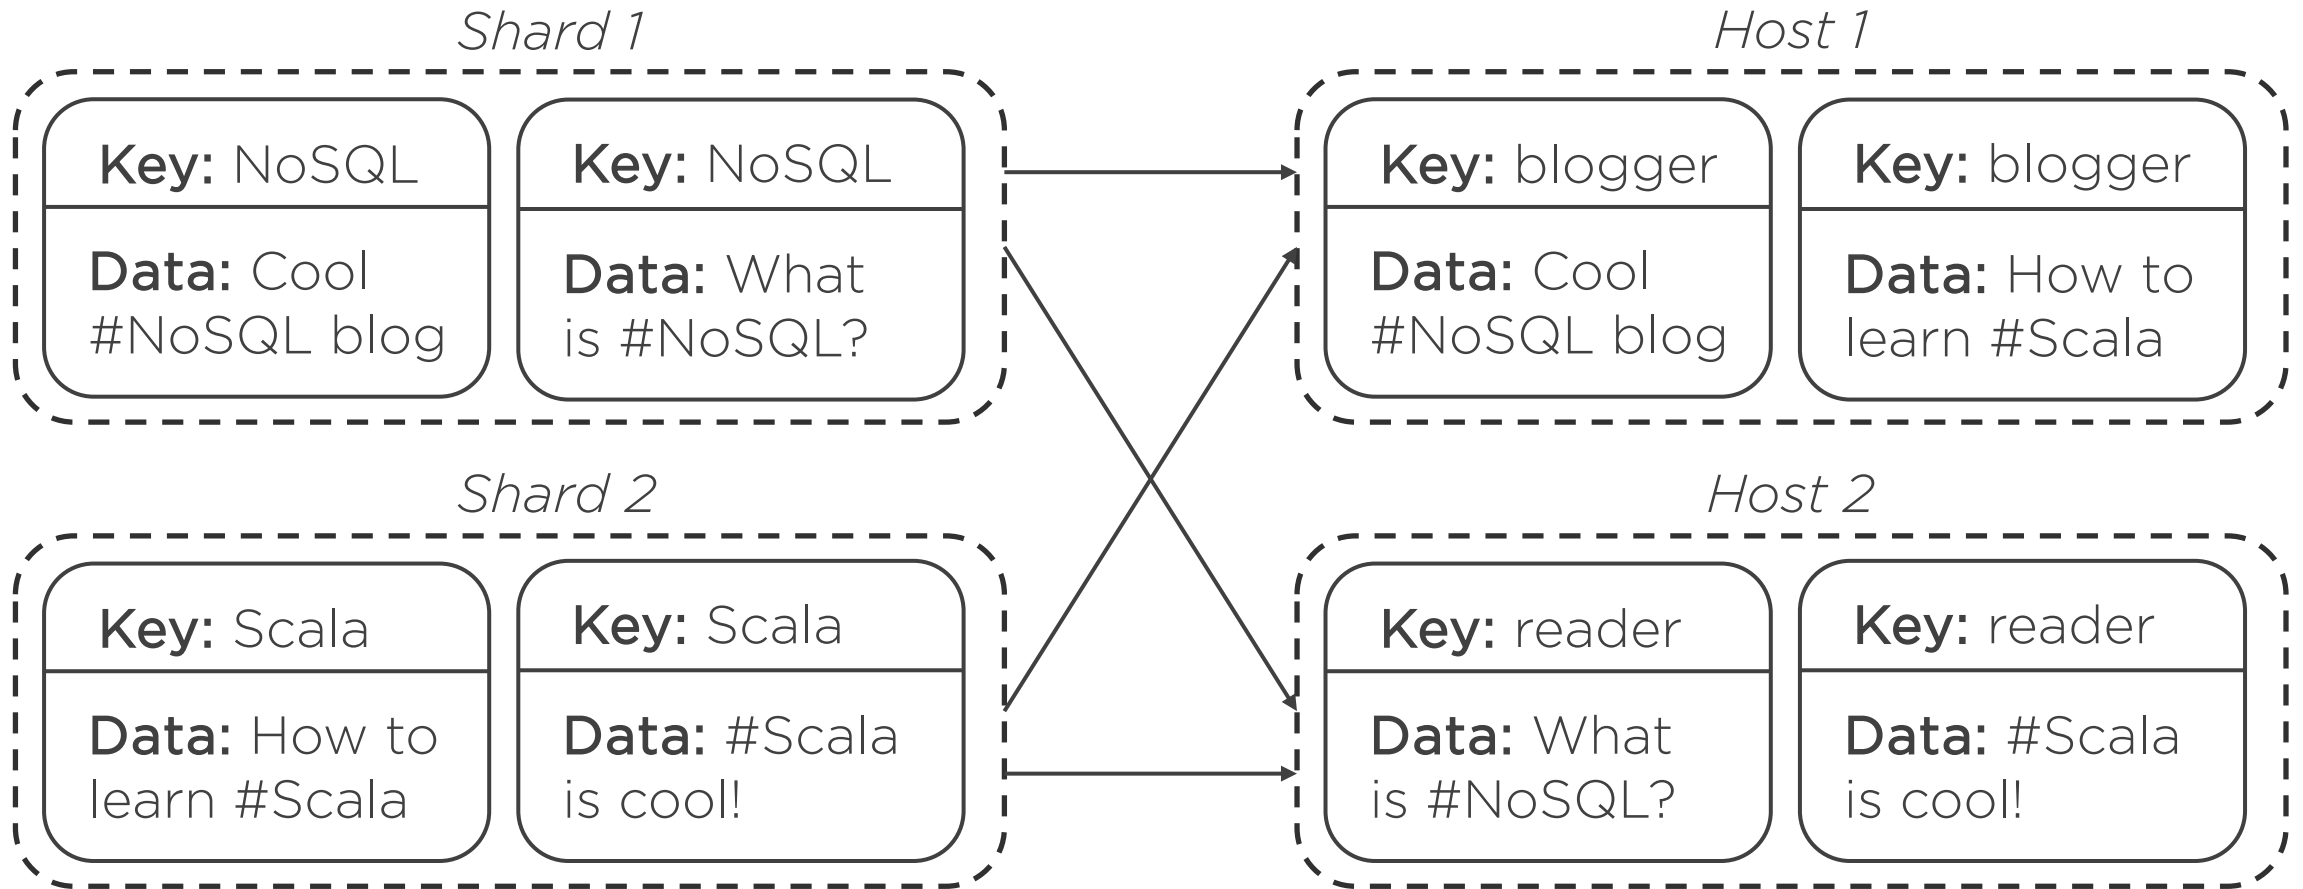
More complex use cases?

Infrastructure maintenance

What Is Missing – Duplicated Records



What Is Missing - Reshuffling



What Is Missing – Late Events

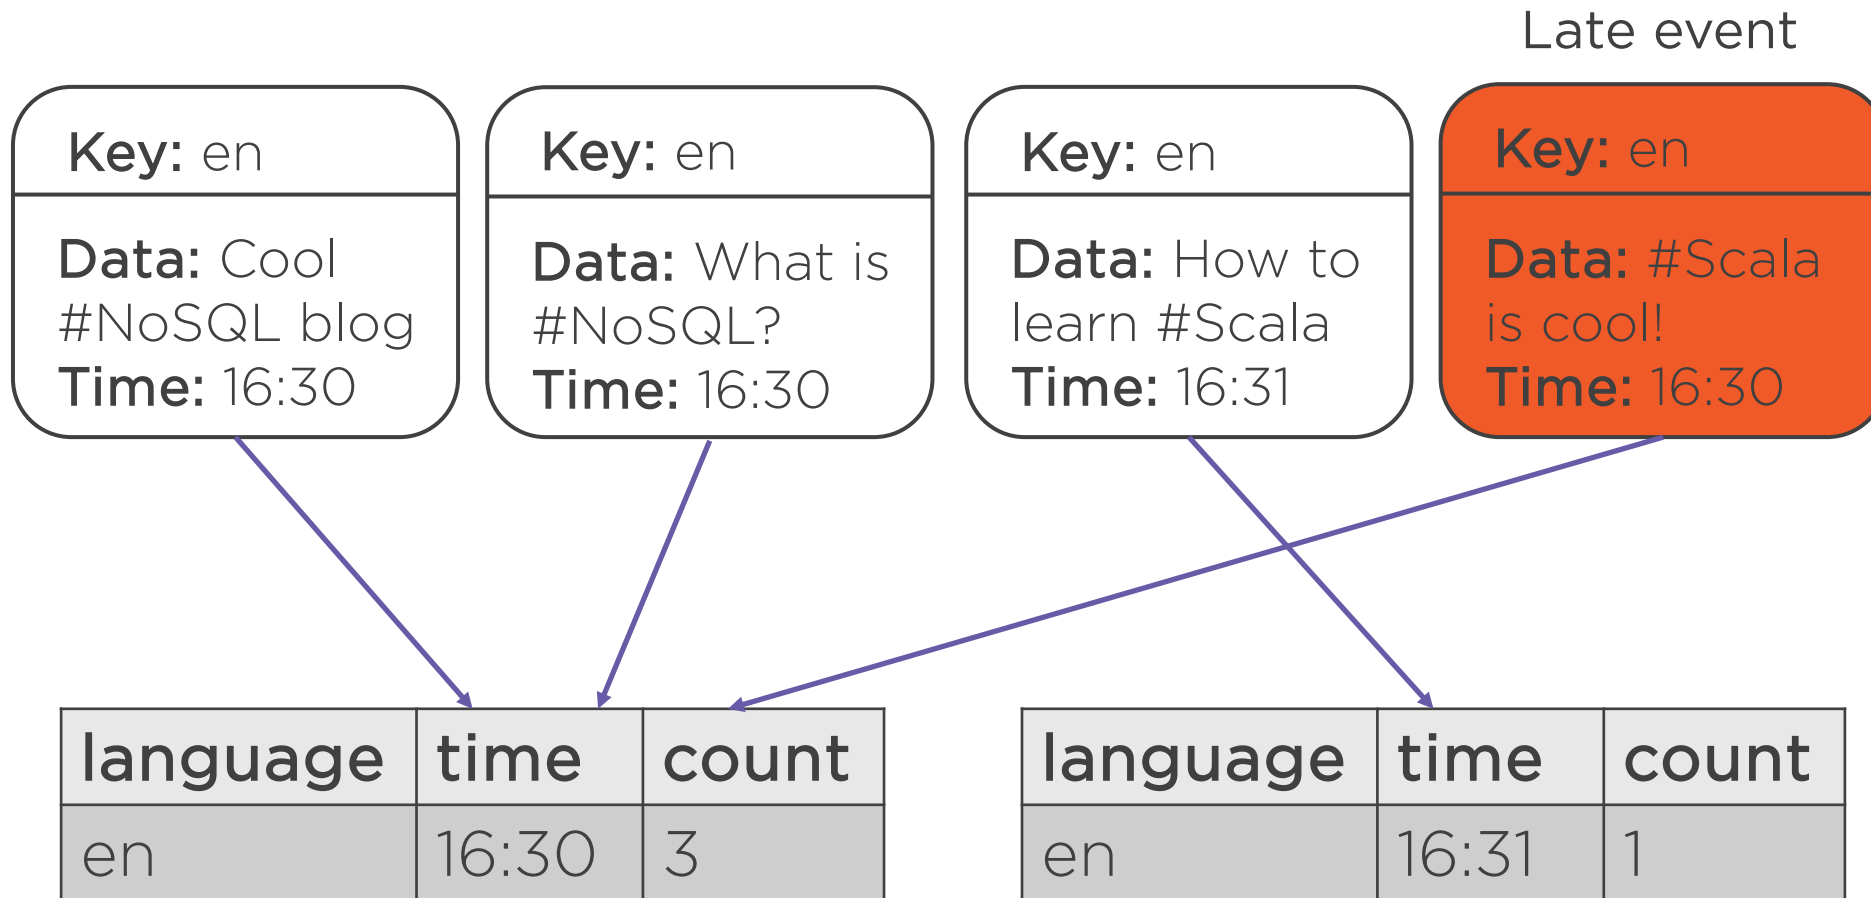


language	time	count
en	16:30	2


language	time	count
en	16:31	2



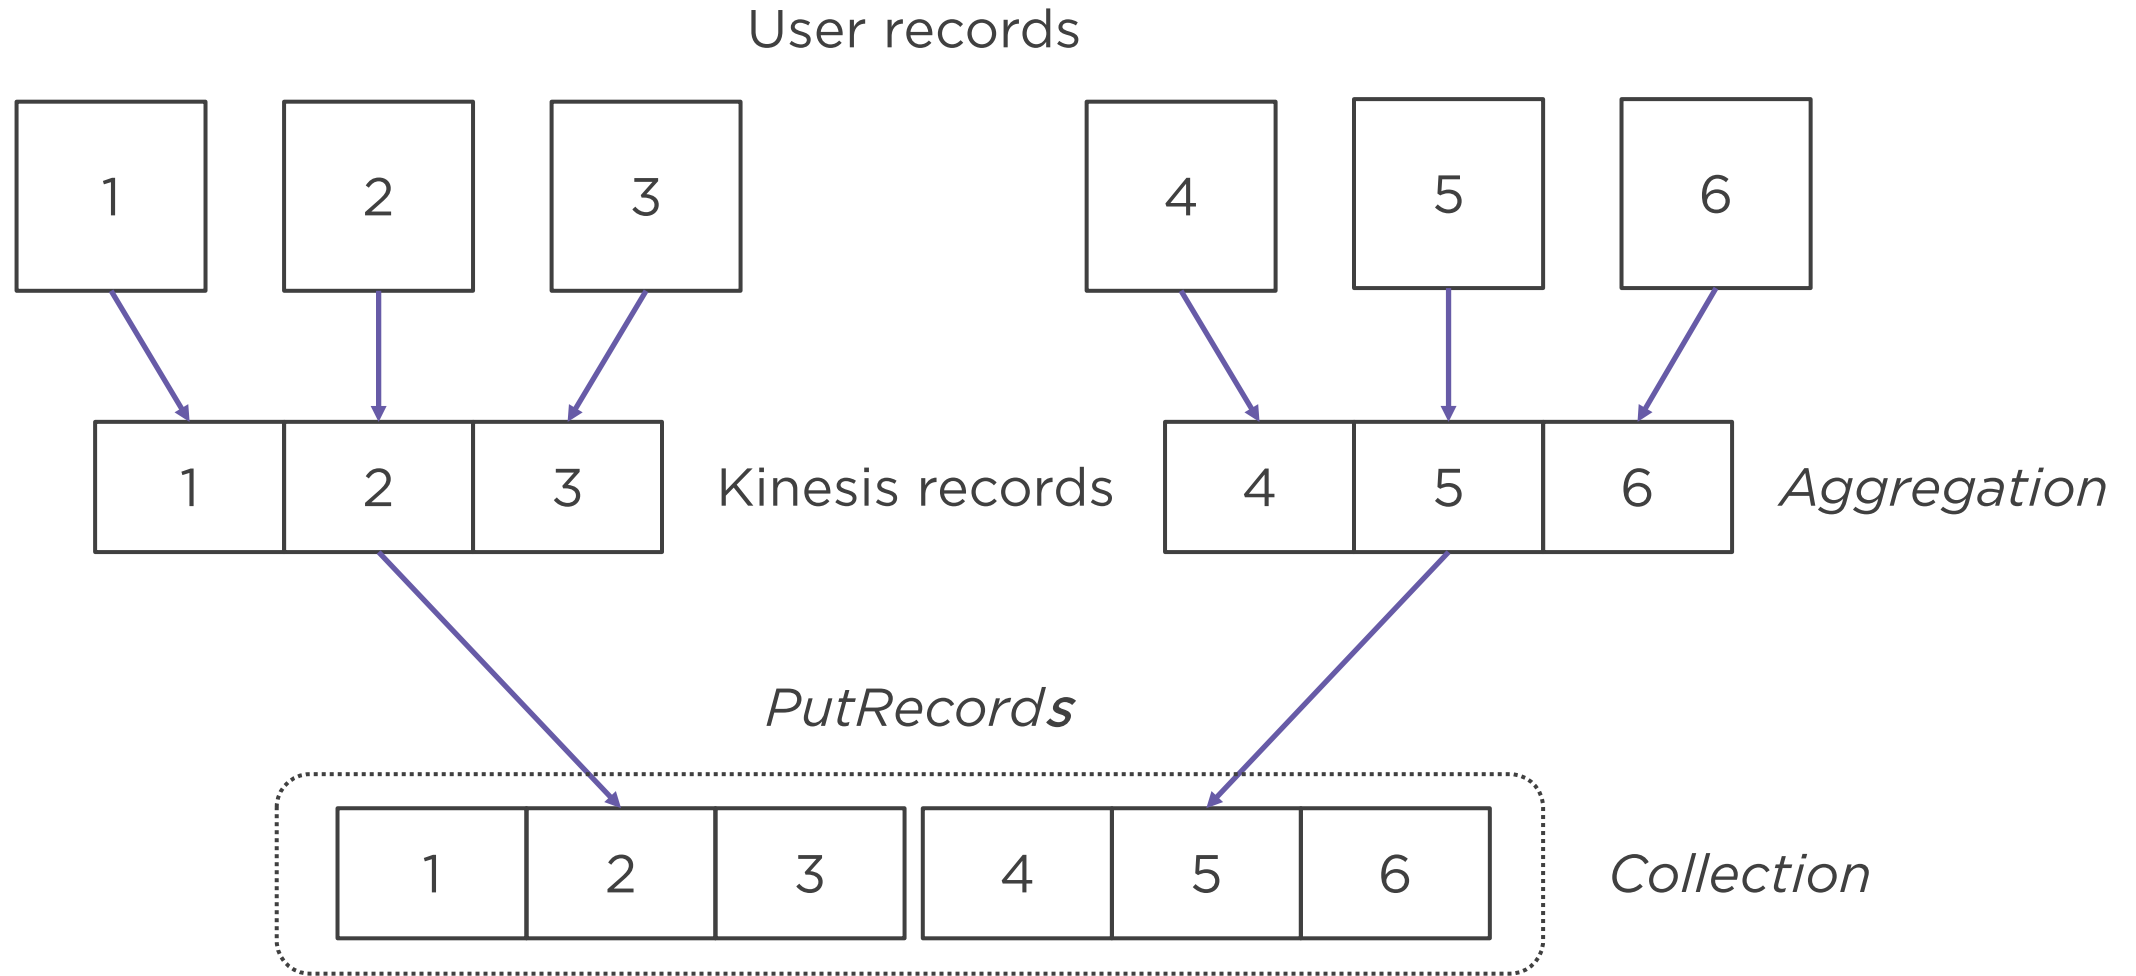
What Is Missing – Late Events



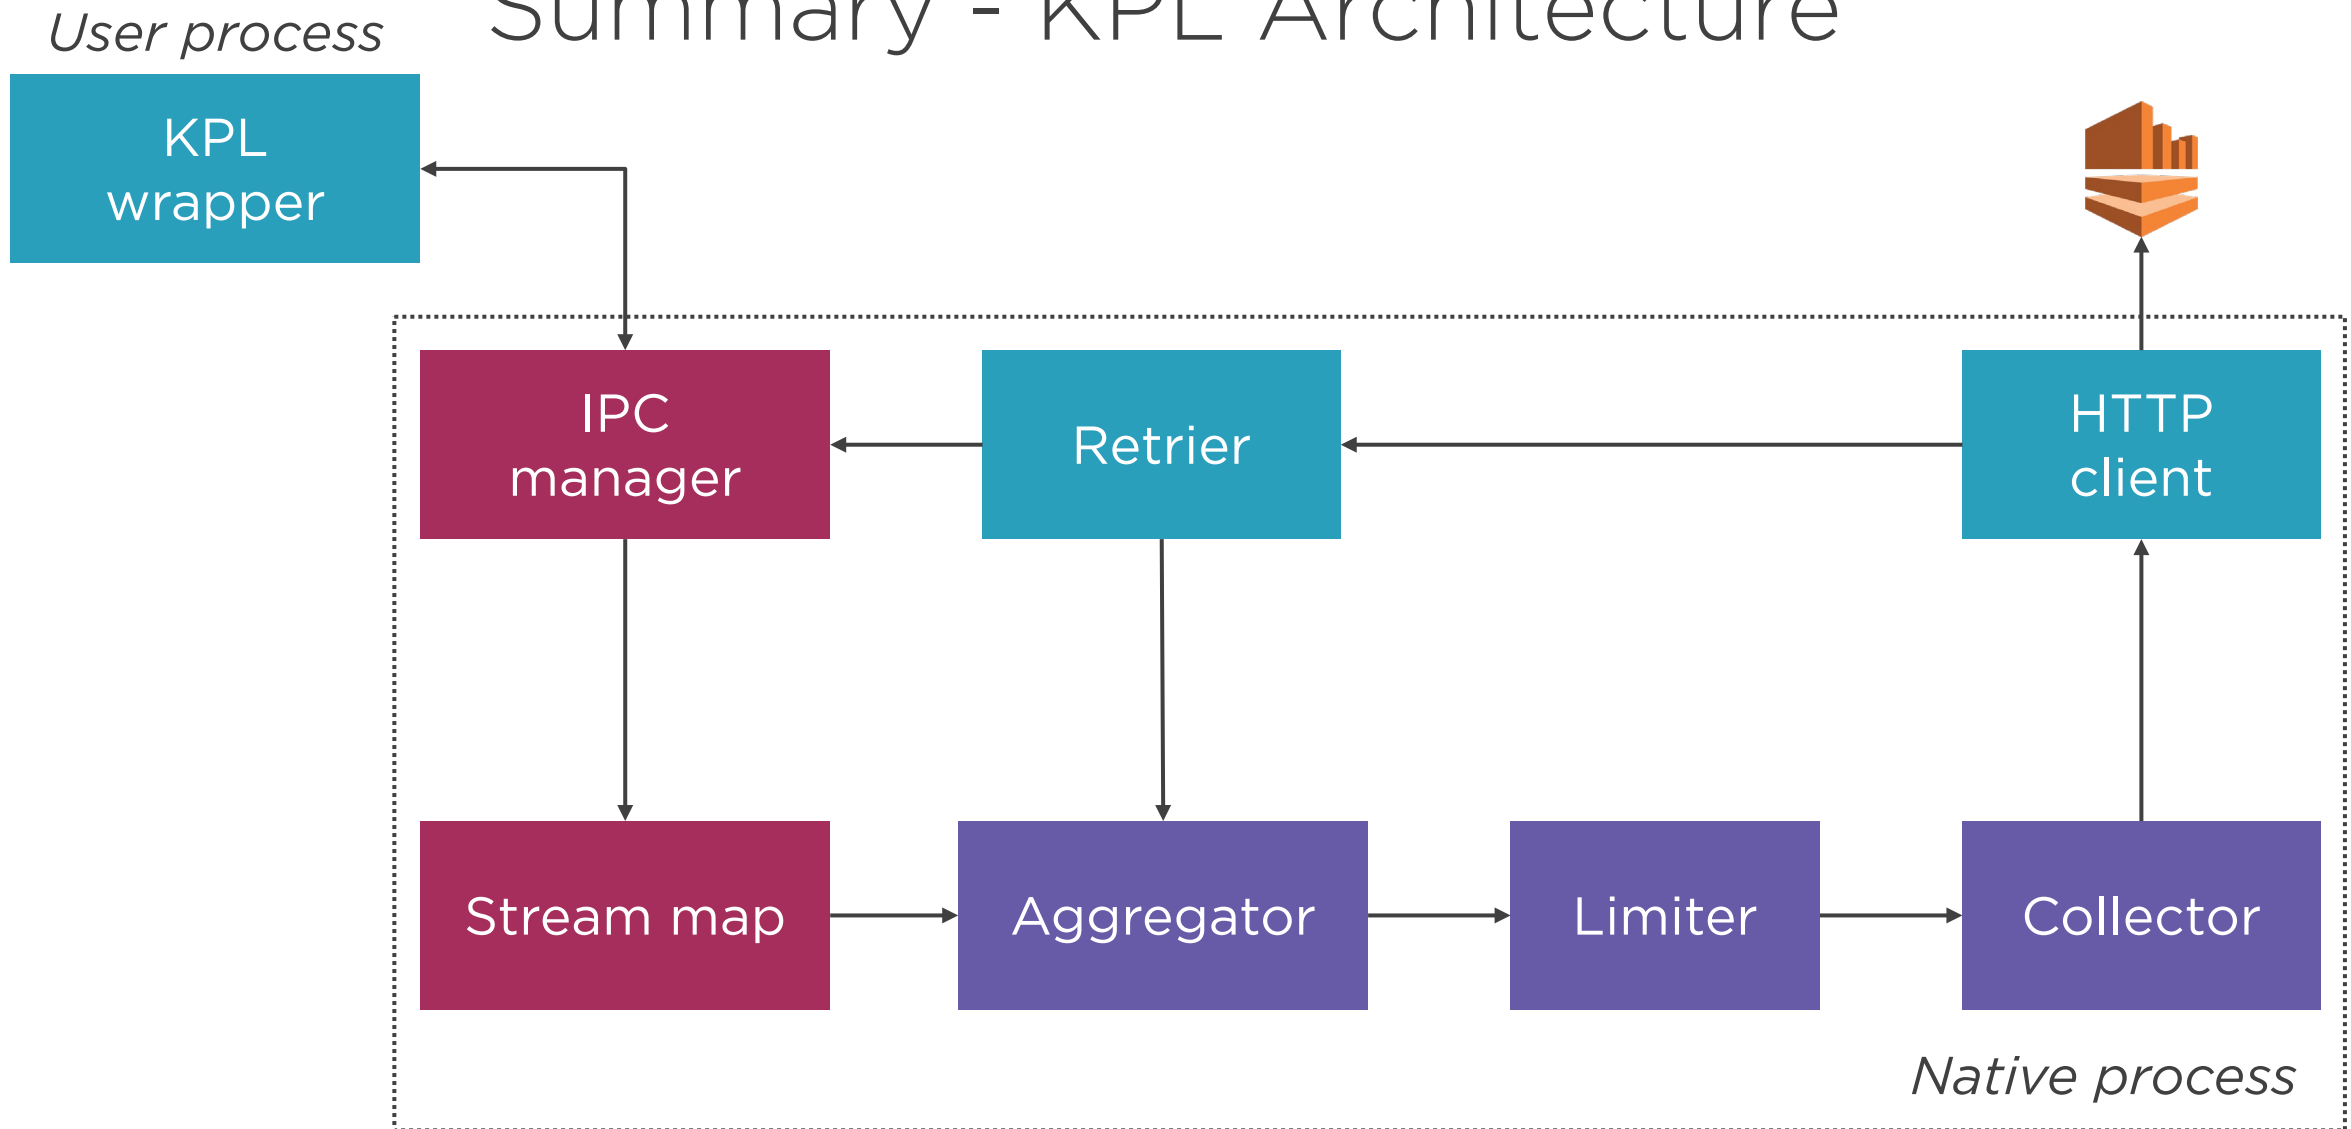
Summary



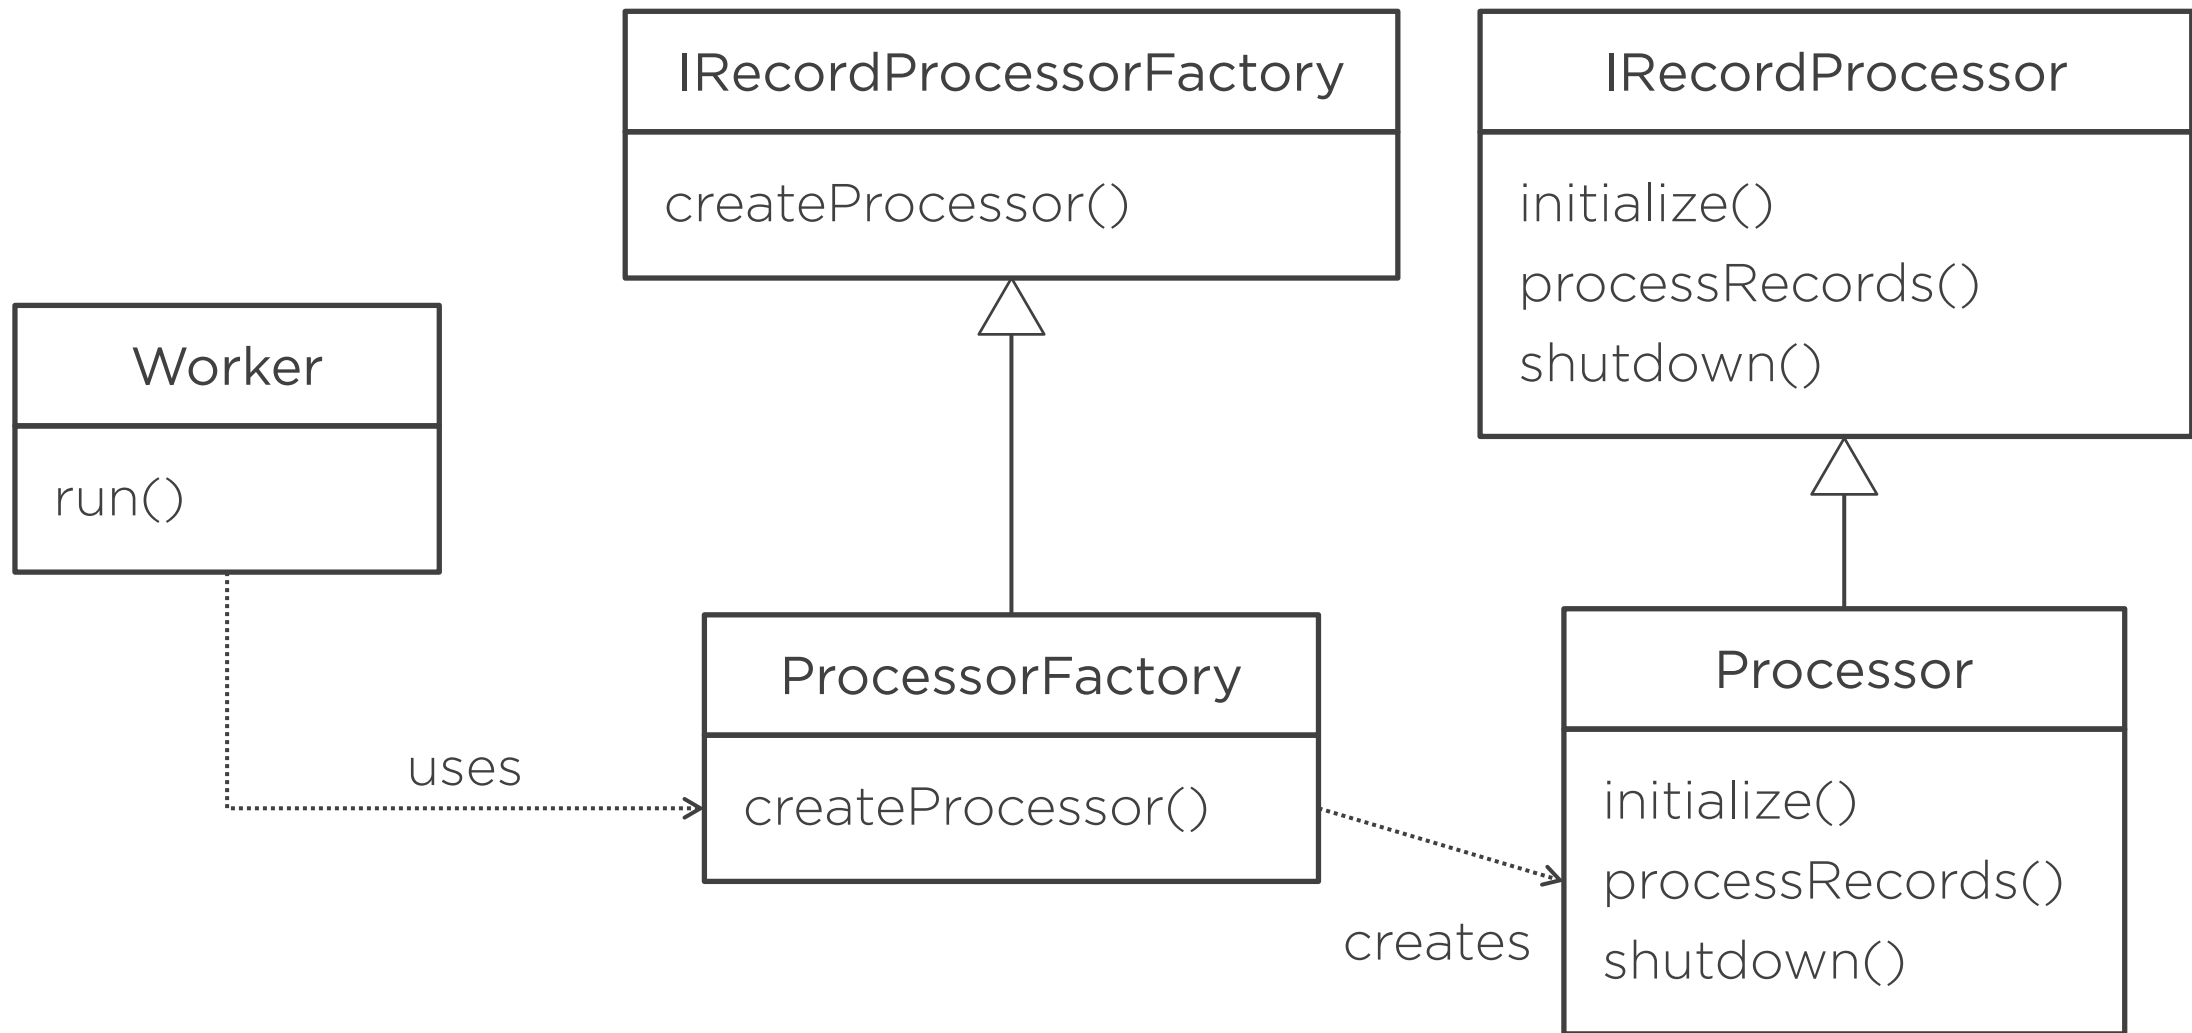
Summary - Aggregation




Summary - KPL Architecture



Summary - Kinesis Client Library



Summary - Scaling with KCL



Summary - Checkpointing

